

Algorithms for Improving the Design and Production of Oligonucleotide Microarrays

Sérgio Anibal de Carvalho Junior

January 2007

PhD thesis submitted to the
Faculty of Technology of Bielefeld University, Germany,
for the degree of Dr. rer. nat.

Supervisor:
Dr. Sven Rahmann

1. Referee: Prof. Dr. Jens Stoye
2. Referee: Prof. Dr. Robert Giegerich

Foreword

Microarrays are a ubiquitous tool in molecular biology with a wide range of applications on a whole-genome scale including high-throughput gene expression analysis, genotyping, and resequencing.

Several different microarray platforms exist, but this thesis focuses on high-density oligonucleotide arrays. The advantage of higher densities arrays is that they allow, for instance, the simultaneous measurement of the expression of several thousand genes at once.

High-density microarrays are usually produced by light-directed combinatorial chemistry that builds the probe sequences base-by-base. Because of the natural properties of light, the quality of a microarray can be improved by carefully designing the physical arrangement, or *layout*, of its probes.

In this thesis, we review models for evaluating the layout of oligonucleotide microarrays and survey algorithmic approaches that can be used in their design.

DNA chips allow to quickly obtain and compare gene expression profiles of different cell or tissue samples. As one of many applications, one hopes to better understand various types and subtypes of cancer and to improve cancer therapy by characterizing the differences between the expression profiles of healthy cells and tumor cells.

The first chapter of this thesis offers an overview of existing microarray technologies and contrasts them with other techniques for gene expression analysis. High-density oligonucleotide arrays are described in detail.

Gene expression analysis with DNA chips is a high-throughput technique that produces massive amounts of data; however, this technique is not error-free. Each step of an experiment must be performed carefully. In particular, the DNA chip must be carefully designed and manufactured. This thesis proposes and describes solutions for some of the algorithmic problems that arise in this phase. These problems are formulated in detail in the last section of Chapter x.

My research started with the general question of how to find oligonucleotide probes for highly homologous transcripts when it cannot be guaranteed that a sufficiently large set of clearly transcript-specific (or *unique*) probes can be found. An interesting future large-scale application could be the individual expression measurement of all splice variations of all genes in the human genome, for example. The basic idea was

to find several *non-unique* probes such that different probes hybridize to different combinations of targets, with the hope that the measured signal can then be decoded into the individual expression levels. Two questions then follow immediately: How does the decoding work, and how can the probes be designed in a way that makes the decoding as simple as possible?

In early 2001, custom probe design was in its infancy. The design of the large commercial chips, such as the Affymetrix GeneChip®, was a well guarded secret of the respective companies. Several other research groups were also beginning to work on probe selection, and their results, made the problem more popular. I soon realized that, before actually working on non-unique probes, more fundamental questions should be addressed. A high-performance large-scale probe selection system for unique probes would be very useful but did not exist. Such a system could then be used as a basis for non-unique probe design. These considerations are reflected in this thesis.

Once a set of probes is chosen for a chip design, the chip must be produced. With several technologies, such as the Affymetrix GeneChip® arrays and febit's **geniom® one** system, the probes are synthesized in situ on the chip with a combination of photolithography and combinatorial chemistry. We consider the problem of optimizing the nucleotide deposition sequence to lower production costs and to decrease the overall error rate during synthesis.

A concluding discussion about the results of this thesis and an outlook into the future can be found in Chapter 10.

Publications. Parts of this thesis have been published in advance.

This thesis also contains previously unpublished material, namely

- new placement algorithm that simultaneously places and re-embeds the probes

Most of the material was published as a book chapter in de Carvalho Jr. and Rahmann (2007).

The stuff on the SCS appeared in de Carvalho Jr. and Rahmann (2005).

Acknowledgments. This work was carried out while I was at the... It has been a very enjoyable experience, thanks to all present and former colleagues and students in the Bioinformatics program.

I would especially like to thank Sven Rahmann for suggesting the topic, providing initial ideas, and for the opportunity to write this thesis under his guidance.

The support of the whole computer service group at the CeBiTec was crucial to the success of this work, and I cannot thank them enough for their help.

I enjoyed fruitful discussions with many people while this thesis was underway; in particular, I am grateful to ... for their valuable hints and comments.

... read early drafts of several chapters of this thesis and helped to improve it in many ways through their comments.

I would like to thank Dr. Peter Hahn and Chris MacPhee who worked on several QAP instances derived from artificial microarray chips for their efforts.

Sérgio A. de Carvalho Jr.

Bielefeld, January 2007

Danke schön!

Danke schön!

Abbreviations and Notation

Abbreviations in alphabetical order

A, A	adenine
Å	Angstrom; $1\text{Å} = 10^{-10}\text{ m} = 1/10\text{ nm}$
aRNA	anti-sense RNA, amplified RNA; complementary to mRNA
bp	base pair(s)
C, C	cytosine
°C	degrees centigrade (Celsius)
cDNA	complementary DNA; DNA synthesized from mRNA by RT
CDS	coding sequence
cRNA	complementary RNA; same as aRNA
CSMS, csms	cumulative statistics of matching statistics
Cy3	Cyanine 3-dNTP; in fluorescently labeled DNA
Cy5	Cyanine 5-dNTP; in fluorescently labeled DNA
DMD	digital micromirror device
DNA	deoxyribonucleic acid
dN	any deoxynucleotide; any of dA, dC, dG, dT
dNTP	any deoxynucleotide-triphosphate; any of dATP, dCTP, dGTP, dTTP
EST	expressed sequence tag
G, G	guanine
HPLC	high-pressure liquid chromatography
IVT	in-vitro transcription
K	Kelvin; physical unit of temperature
LCF	longest common factor
LCFS	longest common factor statistics
M	physical unit of molar concentration: $1\text{ M} = 1\text{ mol/l}$
MCMC	Markov chain Monte Carlo
MES	2-(N-morpholino)ethanesulfonic acid
mol	physical unit of amount of substance $1\text{ mol} = \text{as many entities as atoms in } 0.012\text{ kg of carbon } 12$
mRNA	messenger RNA
MS, ms	matching statistics
N, N	any of the bases A, C, G, T/U
[Na ⁺]	salt (sodium ion) concentration
NaCl	salt (sodium chloride)
NN	nearest neighbor
nt	nucleotide(s)

oligo-(dT)	primer of 12–20 dTs binding to the poly(A) tail of eukaryotic mRNA
P	phosphor
^{33}P	a radioactive phosphor isotope
PAGE	polyacrylamide gel electrophoresis
PCR	polymerase chain reaction
PEM	percent of expected mutations (unit of evolutionary time)
RNA	ribonucleic acid
RT	reverse transcription
RTase	reverse transcriptase; an enzyme
RT-PCR	reverse transcription-polymerase chain reaction
SAGE	serial analysis of gene expression
SAPE	streptavidin phycoerythrin
SBH	sequencing by hybridization
SCS	shortest common supersequence
SCSP	shortest common supersequence problem
SNP	single nucleotide polymorphism
SSC	saline sodium citrate buffer
SSPE	saline sodium phosphate buffer
SVD	singular value decomposition
T, T	thymine
T7	a bacteriophage that infects <i>E. coli</i> ; these viral parasites are numbered T1 through T7
<i>Taq</i>	<i>Thermus aquaticus</i> ; a heat-resistant bacterium
TIFF	tagged image file format
transfrag	transcript fragment
tRNA	transfer RNA
U, U	uracil
UTR	untranslated region; part of mRNA transcripts

Notation by topic

Notation introduced in early chapters is also used in subsequent chapters.

Chapter 1

p	probe sequence; a string
$ p $	length of p
t	transcript or target sequence
T	transcriptome $T = (t_1, \dots, t_n)$
m	number of probes
i	probe index
n	number of transcripts
j	transcript index
θ	hybridization conditions; parameters
$a(p, t; \theta)$	affinity coefficient between p and t , given θ

$p \triangleleft t$	p is a factor (substring) of t
ε	lower affinity limit for specific hybridization

Chapter 2

$x = (x_j)$	n -vector of transcript expression levels
$y = (y_i)$	m -vector of measured probe signals
$A = (A_{ij})$	$m \times n$ -matrix of probe-transcript affinity coefficients; $y = A \cdot x$
$t(i)$	index of the unique target that probe i binds to
a, a_i	particular affinity coefficients
ρ_{ij}	position dependent factor of A_{ij}
β_{ij}	hybridization stability dependent factor of A_{ij}
γ_{ij}	sequence composition dependent factor of A_{ij}
σ_{ij}	self-complementarity dependent factor of A_{ij}
ε_i	stochastic additive noise of probe signal intensity
c_i, c	systematic additive offset of signal intensity
α	scale of signal intensity
η_{ij}	stochastic multiplicative noise of signal intensity
d	differential operator
Δ	denotes a difference
$U; q; w$	internal energy; heat; work
p	pressure
V	volume
T	temperature
$H; H_m^\circ; \Delta_r H^\circ$	enthalpy; standard molar enthalpy; standard reaction enthalpy
$S; S_m^\circ; \Delta_r S^\circ$	entropy; standard molar entropy; standard reaction entropy
$G; G_m^\circ$	Gibbs energy; standard molar Gibbs energy
$\Delta_r G^\circ; \Delta_r G$	standard Gibbs energy of reaction; reaction Gibbs energy
ξ	extent of reaction
R	gas constant; $R = 8.3145 \text{ J K}^{-1} \text{ mol}^{-1}$
K	equilibrium constant
$[S_2]_{\text{eq}}$	equilibrium concentration of reactant S_2 (single stranded probes)
T_M	melting temperature

Chapter 3

$\text{lcf}(p, t)$	length of the longest common factor of p and t
$\text{lcf}^1(p, t)$	same, allowing one mismatch
$\text{lcf}^*(p, t)$	combined measure derived from lcf and lcf^1
$\text{LCF}(p T)$	LCF vector of probe p against transcriptome T
$\text{LCFS}(p T; \Delta)$	LCF statistics of p against T of width Δ
δ	index of LCF statistics, $0 \leq \delta < \Delta$ difference between full probe length and LCF length
$T(i)$	index set of intended targets for probe i
T	transcriptome; see Chapter 1
$U'(p_i T)$	unspecificity of probe i in T ; formal definition
τ	temperature; to avoid confusion with transcriptome T

$\beta(\delta)$	average hybridization probability as a function of δ
ζ	offset constant; $\zeta = \ln(\beta(0)/(1 - \beta(0)))$
$b > 0$	average Gibbs energy per base pair in units of $(R\tau)$
$u(\delta)$	approximation to $\beta(\delta)$
$U(p_i T)$	unspecificity of probe i in T ; practical definition

Chapter 4

$T = (T_c)$	transcriptome of transcript collections T_c ($c = 1..C$)
C	number of collections
$T_c = \{t_{c,j}\}$	transcript collection c with transcripts $t_{c,j}$ ($j = 1..N_c$)
N_c	number of transcripts in collection c
s	target sequence
Σ	alphabet; here $\Sigma = \{\mathbf{A}, \mathbf{C}, \mathbf{G}, \mathbf{T}\}$
$\Sigma^*; \Sigma^+$	all strings (non-empty strings) over alphabet Σ
$\$$	string end marker and separator
\mathbf{X}	unspecified or wildcard character
pos	suffix array
$\text{lcp}(s, t)$	length of the longest common prefix of s and t
lcp	longest common prefix array
q	bucket prefix length
bck	bucket array
$\gamma = \langle Q \rangle$	numerical radix- q representation of a q -gram Q
cl	collection number array
$\text{ms}^{s t} = (\text{ms}_i^{s t})$	matching statistics of s against t
$\text{ms}^{s t;f}$	matching statistics allowing f mismatches of s against t
R_{\min}^0	minimum value of relevant matching statistics
R_{\min}^1	minimum value of relevant ms. with one mismatch
R_{\max}	maximum value of relevant matching statistics
$\text{MS}[i][c]$	matching statistics array; $\text{MS}[i][c] = \text{ms}_i^{s T_c}$
(i, J)	jump at position i to level J
n, m	string lengths; $ s = m, t = n$
\mathbb{P}	generic notation for a probability measure
\mathbb{E}	generic notation for an expectation
$\pi = (\pi_c)_{c \in \Sigma}$	character distribution for a random sequence model
p	probability that two random characters match
q	$q := 1 - p$
K_L	random number of occurrences of a random L -gram in a string
ρ_L	probability that a jump to level L occurs at a fixed position
E_L	expected number of jumps to level L in matching statistics
E_L^+	expected number of jumps to level $\geq L$ in matching statistics
L°	center of the distribution of the LCF of two random strings

Chapter 5

d	probe distance from the transcript's 3'-end
-----	---

$p_i = (d_i, \ell_i)$	probe i defined by end distance d_i and length ℓ_i
$f(d); f_i$	position preference factor for distance d ; $f_i := f(d_i)$
h_i	hairpin formation probability for probe i
U_i	unspecificity for probe i
B_i	badness value for probe i
$\mathcal{B}(\delta)$	additional badness for probes clustering within distance δ
$B'_i; \overline{B}'$	modified badness value for probe i ; threshold \overline{B}'
S	selected probe set

Chapter 6

$\mathbb{1}_{\{\cdot\}}$	generic notation for an indicator function
$S = (s_c)$	target sequences for all collections $c = 1, \dots, C$
$H = (H_{ij})$	binary $m \times n$ hybridization matrix of probes vs. targets
D	design, i.e. a selection of probes; $D \subset \{1, \dots, m\}$
$A^D; H^D$	affinity and hybridization matrix restricted to rows from D
\mathcal{M}	minimum target coverage
\mathcal{A}	average target coverage
A^\top	transpose of A
Σ	diagonal matrix with singular values (σ_j) of A
σ_j	j -th largest singular value
$\text{diag}(\cdot)$	diagonal matrix with specified entries
A^-	pseudo-inverse of A
$\ \cdot\ _2$	Euclidean norm of a vector; spectral norm of a matrix
$\ \cdot\ _\infty$	maximum norm of a vector
$\text{cond}(A)$	condition number of A
$\text{cond}_s(A)$	Skeel condition number of A
μ	maximal number of probes that may be selected
S, T	sets of target indices
$S \Delta T$	symmetric set difference of S and T
$P(S)$	set of probes that hybridize to any target in S
$T(i)$	set of targets that hybridize to probe i
γ	desired minimum target coverage
σ	desired minimum target set separation
$\delta \in \{0, 1\}^m$	binary vector representation of design D
h_j	j -th column of binary hybridization matrix H
$z^{S,T} = (z_i^{S,T})$	indicates whether probe i separates sets S and T
δ^+, h_j^+, z^+	vectors extended for virtual unique probes
B	maximal size of target sets to be separated
$f_0; f_1$	false negative (positive) error rate
$\mathbb{P}(x)$	prior probability of binary expression vector $x \in \{0, 1\}^n$
$\mathbb{P}(y x)$	likelihood of probe signals $y \in \{0, 1\}^m$, given x
$\pi(x) = \mathbb{P}(x y)$	posterior probability of x for fixed y
$q(z x)$	proposal probability for new solution z at current solution x
$\mathcal{N}(x)$	neighborhood of x

$\alpha(x, z)$	acceptance probability for z at x
Q	rate matrix of an evolutionary Markov process
t	evolutionary time
P^t	Markov transition kernel for time t
$\text{expm}(Q)$	matrix exponential; $\text{expm}(Q) := \sum_{n \geq 0} Q^n / n!$

Chapter 7

$\text{csms}^{s t}$	cumulative statistics of matching statistics of s against t
$\text{csms}_{i,\mu}^{s t}$	$\text{csms}_{i,\mu}^{s t}$ refers to position i in s and matches of length $\geq \mu$
$\text{CSMS}[i][\mu]$	CSMS array; $\text{CSMS}[i][\mu] = \text{csms}_{i,\mu}^{s t}$ if $\mu \in [R_{\min}^0, R_{\max}]$
$\sigma_{k,u}^i$	unexpected CSMS at offset k for a probe starting at i
$L_{i,\delta}$	whole genome surrogate for $\text{LCFS}(p_i)_\delta$
L	length of hypothetical transfrags
D	distance between start points of hypothetical transfrags
K	transfrags known to be present
P_1	probes expected to show a signal because of K
M	transfrags of unknown status; $M = \{1..n\} \setminus K$
N	transfrags most likely not present
J	transfrags whose status is inferred by sampling; $J = \{1..n\} \setminus (K \cup N)$
χ	fraction of transfrags hybridizing to additional probes

Chapter 8

Σ	DNA alphabet
π	permutation of the nucleotides
\mathcal{S}	sequence set
e	binary embedding vector of a string into a supersequence
E	embedding matrix with rows e_i , $i = 1.. \mathcal{S} $
\mathcal{U}	upper bound on the supersequence length
\mathcal{L}	lower bound on the supersequence length
L_i	length of the i -th sequence in \mathcal{S}
$N_i(x)$	number of occurrences of $x \in \Sigma$ in the i -th sequence
$N(x)$	$\max_i N_i(x)$
C_i	completion step of the i -th sequence
$W_{i,k}$	number of unproductive steps between $(k-1)$ -st and k -th productive step for sequence i
Φ	cumulative distribution function of the standard Normal distribution

Trademark Notice. Affymetrix® and GeneChip® are registered trademarks of Affymetrix, Inc., Santa Clara, CA, U.S.A. The terms **geniom® one** and **DNA processor®** are registered trademarks of febit AG, Mannheim, Germany.

Contents

Foreword	i
Danke schön!	v
Abbreviations and Notation	vii
1 Introduction	1
1.1 DNA and Genes	1
1.2 Functional Genomics	1
1.3 High-density Oligonucleotide Microarrays	1
2 The Microarray Layout Problem	5
2.1 Problem statement	7
2.2 Border length	7
2.3 Conflict index	9
2.4 Chip quality measures	12
2.5 How hard is the Microarray Layout Problem?	12
3 Placement Algorithms	15
3.1 Optimal masks for uniform arrays	15
3.2 TSP and threading algorithms	16
3.3 Epitaxial placement	18
3.4 Sliding-Window Matching	19
3.5 Row-Epitaxial	20
3.6 Greedy	21
4 MLP and the Quadratic Assignment Problem	25
4.1 Quadratic assignment problem	25
4.2 QAP formulation of the MLP	26
4.3 QAP heuristics	28
4.4 Results	30
4.5 Discussion	32
5 Re-embedding Algorithms	35
5.1 Optimum Single Probe Embedding	35
5.2 Re-embedding algorithms	37

6	Partitioning Algorithms	39
6.1	1-Dimensional Partitioning	40
6.2	2-Dimensional Partitioning	41
6.3	Centroid-based Quadrissection	44
6.4	Pivot Partitioning: Merging partitioning and re-embedding	45
7	Merging Placement and Re-embedding	47
7.1	Greedy+	47
7.2	Results	48
8	Analysis of Affymetrix Microarrays	51
9	Shortest Common Supersequence	53
10	Discussion	55

Chapter 1

Introduction

1.1 DNA and Genes

Choose another title for this sub-section.

1.2 Functional Genomics

Choose another title for this sub-section.

1.3 High-density Oligonucleotide Microarrays

Oligonucleotide microarrays consist of short DNA fragments, called *probes*, affixed or synthesized at specific locations, called *features* or *spots*, of a solid surface. Microarrays are based on the principle of Watson-Crick base pairing. Each probe is a single-stranded DNA molecule of 10 to 70 nucleotides that perfectly matches with a specific part of a *target* molecule. The probes are used to verify whether (or in which quantity) the targets are present in a given biological sample.

This type of microarray was originally designed in the late 1980s as a tool for DNA sequencing, a technology that is known as DNA Sequencing by Hybridization.

TODO: mention uniform arrays.

The first step of a microarray experiment consists of collecting mRNAs or genomic DNA from the cells or tissue under investigation. The mixture to be analyzed is prepared with fluorescent tags and loaded on the array, allowing the targets to hybridize with the probes. Any unbound molecule is washed away, leaving on the array only those molecules that have found a complementary probe. Finally, the array is exposed

to a light source that induces fluorescence, and an optical scanner reads the intensity of light emitted at each spot.

Under ideal conditions, each probe will hybridize only to its target. Thus, it is possible to infer whether a given molecule is present in the sample by checking whether there is light coming from the corresponding spot of the array. The expression level of a gene in a cell can also be inferred because each spot contains several million identical probes, and the strength of the fluorescent signal on a spot is expected to be proportional to the concentration of the target in the sample. In practice, each target is queried by several probes (its *probe set*), and complex statistical calculations are performed to infer the concentration from the observed signals.

High-density microarrays, also called microarray chips, can have more than a million spots, and are thus able to query tens of thousands of genes, covering the entire genome of an organism. The pioneering Affymetrix GeneChip® arrays, for instance, have up to 1.3 million spots on a coated quartz substrate measuring a little over 1 cm². The spots are as narrow as 5 μm (5 microns, or 0.005 mm), and are arranged in a regularly-spaced rectangular grid.

1.3.1 Microarray Production with Photolithographic Masks

GeneChip arrays are produced by combinatorial chemistry and techniques derived from micro-electronics and integrated circuits fabrication. Probes are typically 25 bases long and are synthesized on the chip, in parallel, in a series of repetitive steps. Each step appends the same kind of nucleotide to probes of selected regions of the chip. The selection of which probes receive the nucleotide is achieved by a process called *photolithography* (Fodor et al., 1991, 1993).

Figure 1.1 illustrates this process: The quartz wafer of a GeneChip array is initially coated with a chemical compound topped with a light-sensitive protecting group that is removed when exposed to ultraviolet light, activating the compound for chemical coupling. A lithographic mask is used to direct light and remove the protecting groups of only those positions that should receive the nucleotide of a particular synthesis step. A solution containing adenine (A), thymine (T), cytosine (C) or guanine (G) is then flushed over the chip surface, but the chemical coupling occurs only in those positions that have been previously deprotected. Each coupled nucleotide also bears another protecting group so that the process can be repeated until all probes have been fully synthesized.




Figure 1.1: Affymetrix's probe synthesis via photolithographic masks. The chip is coated with a chemical compound and a light-sensitive protecting group; masks are used to direct light and activate selected probes for chemical coupling; nucleotides are appended to deprotected probes; the process is repeated until all probes have been fully synthesized.

1.3.2 Maskless Microarray Production

Photolithographic masks are notoriously expensive and cannot be changed once they have been manufactured. Thus, any change in the chip layout requires the production of a new set of masks. A similar method of *in situ* synthesis known as Maskless Array Synthesizer (MAS) was later developed to eliminate the need of such masks (Singh-Gasson et al., 1999). Probes are still built by repeating cycles of deprotection and chemical coupling of nucleotides. The illumination, however, relies on an array of miniature mirrors that can be independently controlled to direct or deflect the incidence of light on the chip.

NimbleGen Systems, Inc. uses its own Digital Micromirror Device (DMD) that can control up to 786 000 individual mirrors to produce microarrays with spots as small as $16 \mu\text{m} \times 16 \mu\text{m}$. The Geniom[®] system of febit biotech GmbH, a platform for customized microarray production, also uses a micromirror array to direct the synthesis process.

1.3.3 The Problem of Unintended Illumination

Regardless of which method is used to direct light (masks or micromirror arrays), it is possible that some probes are accidentally activated for chemical coupling because of light diffraction, scattering or internal reflection on the chip surface. This unwanted illumination of regions introduces unexpected nucleotides that change probe sequences, significantly reducing their chances of successful hybridization with their targets. Moreover, these faulty probes may also introduce cross-hybridizations, which can interfere in the experiments performed with the chip.

This problem is more likely to occur near the borders between a masked and an unmasked spot (in the case of maskless synthesis, between a spot that is receiving light and a spot that is not). This observation has given rise to the term *border conflict*.

It turns out that by carefully designing the *arrangement* of the probes on the chip and their *embeddings* (the sequences of masked and unmasked steps used to synthesize each probe), it is possible to reduce the risk of unintended illumination. This issue becomes even more important as there is a need to accommodate more probes on a single chip, which requires the production of spots at higher densities and, consequently, with reduced distances between probes.

In this thesis, we address the problem of designing the layout of a microarray with the goal of reducing the chances of unintended illumination, which we call Microarray Layout Problem (MLP). We use the term *layout* to refer to where and how the probes are synthesized on the chip (their arrangement and their embeddings).

Chapter 2

The Microarray Layout Problem

In this chapter we give a more precise definition of the Microarray Layout Problem (MLP) and define criteria for evaluating a given layout. The description that follows assumes that probes are synthesized with photolithographic masks, but the concepts also apply to the maskless production (with micromirror arrays). Two evaluation criteria are presented: *border length* and *conflict index*. As shown later, the conflict index model can be seen as a generalization of the border length model.

Formally, we have a set of probes $\mathcal{P} = \{p_1, p_2, \dots, p_n\}$, where each $p_k \in \{A, C, G, T\}^*$ is produced by a series of T synthesis steps. Frequently, but not necessarily, all probes have the same length ℓ . Each synthesis step t uses a mask M_t to induce the addition of a particular nucleotide $N_t \in \{A, C, G, T\}$ to a subset of \mathcal{P} (Fig. 2.1). The *nucleotide deposition sequence* $N = N_1 N_2 \dots N_T$ corresponding to the sequence of nucleotides added at each synthesis step is a supersequence of all $p \in \mathcal{P}$.

A microarray chip consists of a set of spots, or sites, $\mathcal{S} = \{s_1, s_2, \dots, s_m\}$, where each spot s is specified by its coordinates on the chip surface and accommodates a unique probe $p_k \in \mathcal{P}$. Note that we usually refer to s as containing a single probe p_k although, in practice, it contains several million copies of it. Each probe is synthesized at a unique spot, hence there is a one-to-one assignment between probes and spots (if we assume that there are as many spots as probes, i.e., $m = n$). Real microarrays may have complex physical structures but, in this thesis, we assume that the spots are arranged in a rectangular grid with n_r rows and n_c columns. We also assume that probes can be assigned to any spot.

In general, a probe can be *embedded* within N in several ways. An embedding of p_k is a T -tuple $\varepsilon_k = (\varepsilon_{k,1}, \varepsilon_{k,2}, \dots, \varepsilon_{k,T})$ in which $\varepsilon_{k,t} = 1$ if probe p_k receives nucleotide N_t (at step t), and 0 otherwise. In particular, a *left-most embedding* is an embedding in which the bases are added as early as possible (as in ε_1 in Fig. 2.1). Similarly, a *right-most embedding* is an embedding in which the bases are added as late as possible (as in ε_8 in Fig. 2.1).

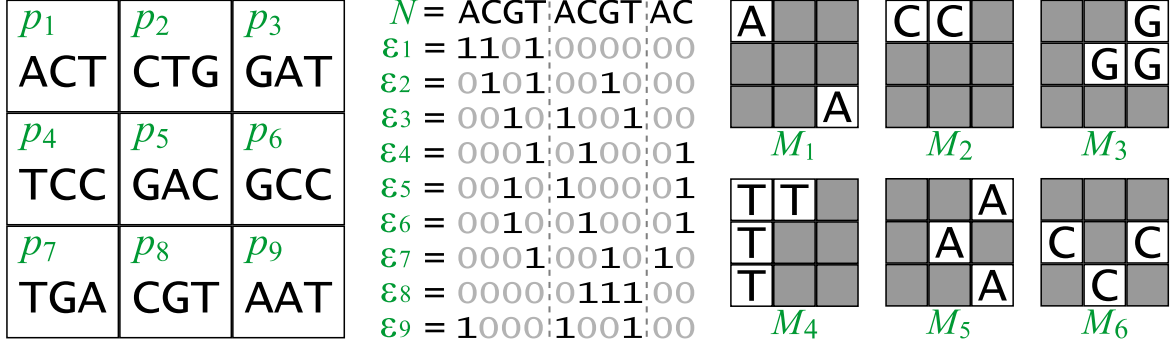


Figure 2.1: Synthesis of a hypothetical 3×3 chip with photolithographic masks. Left: chip layout and the 3-mer probe sequences. Center: deposition sequence with 2.5 cycles (cycles are delimited with dashed lines) and probe embeddings (asynchronous). Right: first six masks (masks 7 to 10 not shown).

We say that an embedding ε_k is *productive* (unmasked) at step t if $\varepsilon_{k,t} = 1$, or *unproductive* (masked) otherwise. The terms productive and unproductive can also be used to denote unmasked and masked spots, respectively.

The deposition sequence is often a repeated permutation of the alphabet, mainly because of its regular structure and because such sequences maximize the number of distinct subsequences (Chase, 1976). The deposition sequence shown in Fig. 2.1 is a 2.5-time repetition of ACGT, and we thus say that it has two and a half *cycles*.

For cyclic deposition sequences, it is possible to distinguish between two types of embeddings: *synchronous* and *asynchronous*. In the former, each probe has exactly one nucleotide synthesized in every cycle of the deposition sequence; hence, 25 cycles or 100 steps are needed to synthesize probes of length 25. In the latter, probes can have any number of nucleotides synthesized in any given cycle, allowing shorter deposition sequences. For this reason, asynchronous embeddings are usually the choice for commercial microarrays. For instance, all GeneChip arrays that we know of can be asynchronously synthesized in 74 steps with 18.5 cycles of TGCA — we refer to this sequence as the *standard Affymetrix deposition sequence* (see Chapter 8).

Ideally, the deposition sequence should be as short as possible in order to reduce manufacturing time, cost and probability of errors (Rahmann, 2003). Finding the shortest deposition sequence to synthesize a set of probes is an instance of a classical computer science problem known as the Shortest Common Supersequence problem, which will be the focus of Chapter 9. For the MLP, however, we assume that N is a fixed sequence given as input.

2.1 Problem statement

Given a set of probes \mathcal{P} , a geometry of spots \mathcal{S} , and a deposition sequence N as specified above, the MLP asks to specify a chip layout (λ, ε) that consists of

1. a bijective assignment $\lambda : \mathcal{S} \rightarrow \{1, \dots, n\}$ that specifies a probe index $k(s)$ for each spot s (meaning that $p_{k(s)}$ will be synthesized at s),
2. an assignment $\varepsilon : \{1, \dots, n\} \rightarrow \{0, 1\}^T$ specifying an embedding $\varepsilon_k = (\varepsilon_{k,1}, \dots, \varepsilon_{k,T})$ for each probe index k , such that $N[\varepsilon_k] := (N_t)_{t:\varepsilon_{k,t}=1} = p_k$,

such that a given penalty function is minimized. We introduce two such penalty functions: total border length and total conflict index.

2.2 Border length

The first formal definition of unintended illumination problem was given by Hannenhalli et al. (2002), who defined the *border length* \mathcal{B}_t of a mask M_t as the number of borders separating masked and unmasked spots at synthesis step t , that is, the number of border conflicts in M_t . Formally,

$$\mathcal{B}_t := \frac{1}{2} \cdot \sum_{s, s' \in \mathcal{S}} \mathbb{1}_{\{s \text{ and } s' \text{ are adjacent}\}} \cdot \mathbb{1}_{\{\varepsilon_{k(s),t} \neq \varepsilon_{k(s'),t}\}}. \quad (2.1)$$

where $\mathbb{1}_{\{cond\}}$ is the indicator function that equals 1 if condition *cond* is true, and 0 otherwise. The *total border length* of a given layout (λ, ε) is the sum of border lengths over all masks, that is

$$\mathcal{B}(\lambda, \varepsilon) := \sum_{t=1}^T \mathcal{B}_t. \quad (2.2)$$

The *Border Length Minimization Problem* was then defined as the problem of finding a layout minimizing the total border length (Hannenhalli et al., 2002). As an example, the six masks shown in Fig. 2.1 have $\mathcal{B}_1 = 4$, $\mathcal{B}_2 = 3$, $\mathcal{B}_3 = 5$, $\mathcal{B}_4 = 4$, $\mathcal{B}_5 = 8$ and $\mathcal{B}_6 = 9$. The total border length of that layout is 52 (masks 7 to 10 not shown).

Hamming distance. In the next chapters, we refer to the *Hamming distance* $H(k, k')$ between the embeddings ε_k and $\varepsilon_{k'}$ as the number of synthesis steps in which they differ. Formally,

$$H(k, k') := \sum_{t=1}^T \mathbb{1}_{\{\varepsilon_{k,t} \neq \varepsilon_{k',t}\}}. \quad (2.3)$$

Note that $H(k, k')$ gives the number of border conflicts generated when probes with embeddings ε_k and $\varepsilon_{k'}$ are placed in adjacent spots.

2.2.1 Lower bounds

Lower bounds for the BLMP with synchronous and asynchronous embeddings were given by Kahng et al. (2002), based on a simple graph formulation. Unfortunately, both lower bounds are not tight, and their computation is time-consuming, especially for large chips.

Synchronous embeddings. Let L be a complete directed graph over the set of probes \mathcal{P} with arcs weighted with the Hamming distance between the (unique) embeddings of the corresponding probes.

Since a probe can have at most four neighbors on the chip, we delete all but the four arcs with the least weights of every node. Furthermore, assuming that the chip is a rectangular grid with n_r rows and n_c columns, we delete the heaviest $2 \cdot (n_r + n_c)$ remaining arcs, because the spots on the borders of the chip have less than four neighbors. It is not difficult to see that the cost of any placement must be greater than the total arc weight of L , and we obtain the following theorem.

Theorem 2.1. *The total arc weight of L is a lower bound on the total border length of the optimum layout with synchronous embeddings.*

Asynchronous embeddings. With asynchronous embeddings, we can construct a similar complete directed graph L' . For the arc weights, however, it is necessary to estimate the minimum number of border conflicts between the two probes (among all of their possible embeddings).

Kahng et al. (2002) observed that the number of bases of a probe p_k that can be “aligned” with bases of $p_{k'}$ cannot exceed the length of $LCS(p_k, p_{k'})$, where $LCS(p_k, p_{k'})$ is the *longest common subsequence* of p_k and $p_{k'}$. Therefore, an arc of L' between probes p_k and $p_{k'}$ can be weighted with $\ell - |LCS(p_k, p_{k'})|$, where ℓ is the length of both probe sequences (assuming probes have the same length).

We can then delete all but the four arcs with the least weights of each probe and, subsequently, the heaviest $2 \cdot (n_r + n_c)$ remaining arcs of L' , to obtain the following theorem.

Theorem 2.2. *The total arc weight of L' is a lower bound on the total border length of the optimum layout with asynchronous embeddings.*

2.3 Conflict index

The border length measures the quality of an individual mask or set of masks. With this model, however, it is not possible to know how the border conflicts are distributed among the probes. Ideally, all probes should have roughly the same risk of being damaged by unintended illumination, so that all signals are affected in approximately the same way.

The *conflict index* is a quality measure defined with the aim of estimating the risk of damaging probes at a particular spot (de Carvalho Jr. and Rahmann, 2006b); it is thus a per-spot or per-probe measure instead of a per-mask measure. Additionally, it takes into account two practical considerations observed by Kahng et al. (2003a):

- a) stray light might activate not only adjacent neighbors but also spots that lie as far as three cells away from the targeted spot;
- b) imperfections produced in the middle of a probe are more harmful than in its extremities.

For a proposed layout (k, ε) , the conflict index $\mathcal{C}(s)$ of a spot s whose probe $p_{k(s)}$ is synthesized in T masking steps according to its embedding vector $\varepsilon_{k(s)}$ is

$$\mathcal{C}(s) := \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_{k(s),t}=0\}} \cdot \omega(\varepsilon_{k(s),t}) \cdot \sum_{\substack{s': \text{neighbor} \\ \text{of } s}} \mathbb{1}_{\{\varepsilon_{k(s'),t}=1\}} \cdot \gamma(s, s') \right). \quad (2.4)$$

The indicator functions ensure the following conflict condition: During step t , there is a conflict at spot s if and only if s is masked ($\varepsilon_{k(s),t} = 0$) and a close neighbor s' is unmasked ($\varepsilon_{k(s'),t} = 1$) — since light directed at s' may somehow reach s . When s is unmasked, it does not matter if it accidentally receives light targeted at a neighbor, and when s' is masked, there is no risk that it damages probes of s since it is not receiving light.

Function $\gamma(s, s')$ is a “closeness” measure between s and s' (to account for observation a). We defined it as

$$\gamma(s, s') := (d(s, s'))^{-2}, \quad (2.5)$$

where $d(s, s')$ is the Euclidean distance between the spots s and s' . Note that in (2.4), s' ranges over all neighboring spots that are at most three cells away from s (see Fig. 2.2, left), which is in accordance with observation a.

In general, we use the terms *close neighbor* or simply *neighbor* of a spot s to refer to a spot s' that is at most three cells away (vertically and horizontally) from s . In other words, s' is inside a 7×7 region centered around s . This is in contrast to the terms *direct* or *immediate neighbor* of s , used to denote a spot s' that is adjacent to




Figure 2.2: Ranges of values for both γ and ω on a typical Affymetrix chip where probes of length $\ell = 25$ are synthesized in $T = 74$ masking steps. Left: approximate values of the distance-dependent weighting function $\gamma(s, s')$ for a spot s in the center and close neighbors s' . Right: position-dependent weights $\omega(\varepsilon, t)$ on the y-axis for each value of $b_{\varepsilon, t} \in \{0, \dots, 25\}$ on the x-axis, using $\theta = 5/\ell$ and $c = 1/\exp(\theta)$.

s (in other words, when s' shares a common border with s on the chip). Obviously, an immediate neighbor s' is also a close neighbor of s .

The position-dependent weighting function $\omega(\varepsilon, t)$ accounts for the significance of the location inside the probe where the undesired nucleotide is introduced in case of accidental illumination (observation b). We defined it as:

$$\omega(\varepsilon, t) := c \cdot \exp(\theta \cdot \lambda(\varepsilon, t)) \quad (2.6)$$

where $c > 0$ and $\theta > 0$ are constants, and for $1 \leq t \leq T$,

$$\lambda(\varepsilon, t) := 1 + \min(b_{\varepsilon, t}, \ell_\varepsilon - b_{\varepsilon, t}), \quad (2.7)$$

$$b_{\varepsilon, t} := \sum_{t'=1}^t \varepsilon_{t'}, \quad \ell_\varepsilon := \sum_{t=1}^T \varepsilon_t = b_{\varepsilon, T}. \quad (2.8)$$

In other words, ℓ_ε is the length of the final probe specified by ε (equal to the number of ones in the embedding), and $b_{\varepsilon, t}$ denotes the number of nucleotides synthesized up to and including step t .

We can now speak of the *total conflict index* of a given layout (λ, ε) as the sum of conflict indices over all spots, that is

$$\mathcal{C}(\lambda, \varepsilon) := \sum_s \mathcal{C}(s). \quad (2.9)$$

Conflict index distance. Many of the algorithms discussed in later chapters were initially developed for border length minimization, and they usually rely on the Hamming distance defined earlier (2.3). We have adapted some of these algorithms to work with conflict index minimization by using the *conflict index distance*, which extends the Hamming distance by taking into account the position inside the probe where the conflict occurs (observation b). The conflict index distance $C(k, k')$ between the embeddings ε_k and $\varepsilon_{k'}$ is defined as:

$$C(k, k') := \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_{k,t}=0 \text{ and } \varepsilon_{k',t}=1\}} \cdot \omega(\varepsilon_k, t) + \mathbb{1}_{\{\varepsilon_{k',t}=0 \text{ and } \varepsilon_{k,t}=1\}} \cdot \omega(\varepsilon_{k'}, t) \right). \quad (2.10)$$

The conflict index distance $C(k, k')$ can be interpreted as the sum of the conflict indices resulting from placing probes with embeddings ε_k and $\varepsilon_{k'}$ at hypothetical neighboring spots, ignoring the distance between these spots (note that there is no dependency on γ) and the conflicts generated by other neighbors.

2.3.1 The choices of γ and ω

The conflict index $\mathcal{C}(s)$ attempts to estimate the risk of damaging the probes of a spot s due to unintended illumination. The definitions of γ and ω given here are an arbitrary choice in an attempt to capture the characteristics of the problem.

However, the most appropriate choice of γ depend on several attributes of the specific technology utilized to produce the chips such as the size of the spots, the density of the probes on the chip, the physical properties of the light being used (intensity, frequency, etc.), the distance between the light source and the mask, and the distance between the mask (or the micromirrors) and the chip surface.

The most appropriate choice of ω depend on the chemical properties of the hybridization between probes and targets. Although it is generally agreed that the chances of a successful hybridization are higher if a mismatched base occurs at the extremities of the formed duplex instead of at its center (Hubbell et al., 1999), the precise effects of this position is not yet fully understood and has been an active topic of research (Binder and Preibisch, 2005).

We propose the use of an exponential function, so that ω grows exponentially from the extremities of the probe to its center (see Fig. 2.2, right). The motivation behind this definition is that the probability of a successful stable hybridization of a probe with its target should increase exponentially with the absolute value of its Gibbs free energy, which increases linearly with the length of the longest perfect match between probe and target.

The parameter θ controls how steeply the exponential weighting function rises towards the middle of the probe. In our experiments, unless stated otherwise, we use probes of length $\ell = 25$, and parameters $\theta = 5/\ell$ and $c = 1/\exp(\theta)$.

Finding the best choice of γ and ω for a particular technology is beyond the scope of this thesis. We note, however, that all algorithms discussed in the next chapters were developed to work independently of the values given by these functions. In other words, should γ and ω be defined differently, no changes to the algorithms are necessary.

2.4 Chip quality measures

Most of the algorithms discussed in the next chapters can work with border length as well as conflict index minimization. The relation between these two measures becomes clear if $\gamma(s, s')$ and $\omega(\varepsilon, t)$ are re-defined as follows: Set $\gamma(s, s') := 1$ if s' is a direct neighbor of s , and $:= 0$ otherwise. Also, set $c = 1/2$ and $\theta = 0$ so that $\omega(\varepsilon, t) := 1/2$ independently of the position in the probe where the conflict occurs. Now $\sum_s \mathcal{C}(s) = \sum_{t=1}^T \mathcal{B}_t$; that is, total border length is equivalent to the total conflict index for a particular choice of γ and ω . For the choices (2.5) and (2.6), they are not equivalent but still correlated, since a good layout has low border lengths as well as low conflict indices.

To better compare border lengths for chips of different sizes, we usually divide the total border length by the number n_b of internal borders of the chip, which equals $n_r(n_c - 1) + n_c(n_r - 1)$ if the chip is a rectangular grid with n_r rows and n_c columns. We thus call $\mathcal{B}(\lambda, \varepsilon)/n_b$ the *normalized border length* of a given layout (λ, ε) . This can be further divided by the number of synthesis steps to give the *normalized border length per mask* $\mathcal{B}(\lambda, \varepsilon)/(n_b \cdot T)$. We may also refer to the normalized border length of a particular mask M_t as B_t/n_b .

Similarly, it is useful to divide the total conflict index by the number of probes on the chip, and we define the *average conflict index* of a layout as $\mathcal{C}(\lambda, \varepsilon)/|\mathcal{P}|$.

2.5 How hard is the Microarray Layout Problem?

The MLP appears to be hard because of the super-exponential number of possible arrangements, although no NP-hardness proof is yet known. A formulation of the MLP as a Quadratic Assignment Problem (QAP) is given in Chapter 4. The QAP is a classical combinatorial optimization problem that is, in general, NP-hard, and particularly hard to solve in practice (Çela, 1997). Optimal solutions are thus unlikely

to be found even for small chips and even if we assume that all probes have a single predefined embedding.

If we consider all possible embeddings (up to several million for a typical Affymetrix probe), the MLP is even harder. For this reason, the problem has been traditionally tackled in two phases. First, an initial embedding of the probes is fixed and an arrangement of these embeddings on the chip with minimum conflicts is sought. This is usually referred to as the *placement* phase. Second, a post-placement optimization phase *re-embeds* the probes considering their location on the chip, in such a way that the conflicts with neighboring spots are further reduced. Often, the chip is *partitioned* into smaller sub-regions before the placement phase in order to reduce running times, especially on larger chips.

The most important placement algorithms are surveyed in Chapter 3, whereas re-embedding algorithms are discussed in Chapter 5. Partitioning algorithms are the focus of Chapter 6. Finally, we present recent developments that simultaneously place and re-embed probes in Chapter 7.

Chapter 3

Placement Algorithms

The input for a placement algorithm consists of a geometry of spots \mathcal{S} , the deposition sequence N , and a set of probes \mathcal{P} , where each probe is assumed to have at least one embedding in N . The output is a one-to-one assignment λ of probes to spots. If there are more spots than probes to place, one can add enough “empty” probes that do not introduce any conflicts with the other probes (since light is never directed to their spots).

All algorithms discussed in this section assume that an initial embedding of the probes is given, which can be a left-most, right-most, synchronous or otherwise pre-computed embedding — a placement algorithm typically does not change the given embeddings.

3.1 Optimal masks for uniform arrays

Feldman and Pevzner (1994) were the first to formally address the unintended illumination problem. They showed how a placement for an *uniform array* with minimum number of border conflicts can be constructed using a two-dimensional Gray code. Uniform arrays are arrays containing all 4^ℓ probes of a given length ℓ , which require a deposition sequence of length $4 \cdot \ell$. These arrays were initially developed for the technique known as Sequencing by Hybridization (SBH).

In general, the term Gray code refers to an ordering of a set of elements in which successive elements differ in some pre-specified, usually small, way (Savage, 1997). The construction of Feldman and Pevzner is based on a two-dimensional Gray code composed of string of length ℓ over a four-letter alphabet. It generates an $2^\ell \times 2^\ell$ array filled with ℓ -mer probes in which each pair of adjacent probes (horizontally or vertically) differs by exactly one letter. This construction is illustrated in Fig. 3.1. An $(\ell + 1)$ -mer array is constructed by first copying the ℓ -mer array into the upper left quadrant of the $(\ell + 1)$ -mer array and reflecting it horizontally and vertically into the other three quadrants. The letter in front of the probes in the upper left quadrant of

A	C	AA	AC	CC	CA	AAA	AAC	ACC	ACA	CCA	CCC	CAC	CAA
T	G	AT	AG	CG	CT	AAT	AAG	ACG	ACT	CCT	CCG	CAG	CAT
		TT	TG	GG	GT	ATT	ATG	AGG	AGT	CGT	CGG	CTG	CTT
		TA	TC	GC	GA	ATA	ATC	AGC	AGA	CGA	CGC	CTC	CTA
						TTA	TTC	TGC	TGA	GGA	GGC	GTC	GTA
						TTT	TTG	TGG	TGT	GGT	GGG	GTG	GTT
						TAT	TAG	TCG	TCT	GCT	GCG	GAG	GAT
						TAA	TAC	TCC	TCA	GCA	GCC	GAC	GAA

Figure 3.1: Construction of a placement for uniform arrays (containing the complete set of ℓ -mer probes) based on a two-dimensional Gray code, resulting in layouts with minimum number of border conflicts.

the ℓ -mer array is added to all probes in the upper left quadrant of the $(\ell + 1)$ -mer array. The probes of the other three quadrants are extended in the same way.

It can be shown that such placement generates masks with a minimum number of border conflicts if the probes are synchronously embedded (see Fig. 3.2). However, because this construction is restricted to uniform arrays and synchronous embeddings, it is of limited practical importance for current microarrays.

3.2 TSP and threading algorithms

The border length problem on arrays of arbitrary probes was first discussed by Hanenhalli et al. (2002). The article reports that the first Affymetrix chips were designed using a heuristic for the traveling salesman problem (TSP). The idea is to build a weighted graph with nodes representing probes, and edges containing the Hamming distances between their embeddings (see Eq.2.3). A TSP tour on this graph is heuristically constructed and *threaded* on the array in a row-by-row fashion (Fig. 3.3a).

For uniform arrays, every solution of the TSP corresponds to a (one-dimensional) Gray code since consecutive elements in the tour differ in only one position, thus minimizing border conflicts between consecutive probes. For general arrays, a TSP solution also reduces border conflicts as consecutive probes in the tour are likely to be similar.




Figure 3.2: Masks for the 8×8 uniform array of Fig. 3.1 when probes are synchronously embedded into $\{ACGT\}^3$. Masked spots are represented by shaded squares, unmasked spots by white squares. Note that masks of the same cycle have the same number of border conflicts.




Figure 3.3: Different ways of *threading* probes on a chip. a) Standard row-by-row (0-threading); b) 1-threading; c) 2-threading.

Threading the tour row-by-row leads to an arrangement where consecutive probes in the same row have few border conflicts, but probes in the same column may have very different embeddings.

Another problem of this approach is that the TSP is known to be NP-hard (Gross and Yellen, 2004), so computing an optimal TSP tour even for a small 300×300 array is not feasible, and only fast approximation algorithms are suitable. In practice, Hannenhalli and co-workers managed to achieve marginal improvements in tour cost using the 2-opt algorithm for TSP of Lin and Kernighan (1973) and an algorithm for weighted matching due to Gabow (1976). Unfortunately, their efforts resulted in only 1.05% reduction in tour cost for a chip with 66 000 probes when compared to the greedy TSP algorithm initially used at Affymetrix.

Since improvements in the cost of the TSP tour seemed unlikely, Hannenhalli and co-workers turned their attention to the threading of the tour on the chip. They studied several threading alternatives, which they collectively called *k-threading* (Fig. 3.3b,c).

A *k*-threading is a variation of the standard row-by-row threading, in which the right-to-left and left-to-right paths are interspaced with alternating upward and downward movements over *k* sites (the row-by-row threading can be seen as a *k*-threading with *k* = 0); *k* is called the *amplitude* of the threading. Hannenhalli and co-workers experimentally observed that 1-threading may reduce total border length of layouts constructed with TSP tours in up to 20% for large chips when compared to row-by-row threading.

3.3 Epitaxial placement

A different strategy inspired by techniques used in the design of VLSI circuits, called Epitaxial placement, was proposed by Kahng et al. (2002). It essentially grows a placement around a single starting “seed” using a greedy heuristic. Although it was originally designed for chips with synchronous embeddings, it can be trivially implemented for asynchronous embeddings as well.

The algorithm starts by placing a random probe in the center of the array and continues to insert probes in spots adjacent to already-filled spots. Priority is given to spots whose all four neighbors are full, in which case a probe with the minimum number of border conflicts with the neighbors is placed. Otherwise, all spots with $i \geq 1$ filled neighbors are examined. For each spot s , the algorithm finds a non-assigned probe p whose number of border conflicts with the filled neighbors of s , $c(s, p)$, is minimal, and assigns a normalized cost $\bar{c}(s, p) := \sigma_i \cdot c(s, p)/i$ for this assignment, where $0 < \sigma_i \leq 1$ are scaling coefficients (the authors propose $\sigma_1 = 1$, $\sigma_2 = 0.8$, and $\sigma_3 = 0.6$). The assignment with minimum $\bar{c}(s, p)$ is made and the procedure is repeated until all probes have been placed.

In order to avoid repeated cost computations, the authors propose keeping a list of probes candidates, for each spot, sorted by their normalized costs. This list must be compiled at most four times, but on the average it is only computer two times.

With this algorithm, Kahng and co-workers claim a further 10% reduction in border conflicts over TSP + 1-threading. However, the Epitaxial algorithm have at least quadratic time complexity as it examines every non-placed probe to fill each spot, and large memory requirements if a list of probe candidates is kept for each spot. Hence, like the TSP approach, it does not scale well to large chips. In their experiments, the Epitaxial algorithm needed 274 seconds to design a 100×100 chip, but 4 441 seconds




Figure 3.4: Sliding-Window Matching algorithm. a) Initial arrangement of probes $p_1 \dots p_{16}$ inside a 4×4 window and the selected independent set of spots (shaded). b) Bipartite graph and a minimum weight perfect matching (dark edges). c) New arrangement inside the window according to the perfect matching.

to design a 200×200 chip, that is, a 16.2-fold increase in running time for a 4-fold increase in number of spots. Chips of larger dimensions could not be computed because of prohibitively large running time and memory requirements.

3.4 Sliding-Window Matching

The Sliding-Window Matching algorithm (Kahng et al., 2003a), SWM for short, is not exactly a placement algorithm as it iteratively improves an existing placement that can be constructed, for instance, by TSP + 1-threading.

The authors noted that the TSP tour can be conveniently substituted by lexicographically sorting the probe sequences or their binary embedding vectors with a linear-time radix sort. The sorting is several times faster but it is also likely to produce a worse initial placement than the TSP, with consecutive embeddings being similar only in their first synthesis steps. The authors argue that this is of little importance in practice given that this placement is only used as a starting point for the SWM algorithm, and the lexicographical sorting should be the choice for large microarrays because computing a TSP tour takes prohibitively long for chips larger than 500×500 .

As its name implies, SWM works inside a window that starts at the top left of the chip and slides from left to right, top to bottom, while maintaining a certain amount of overlap between each iteration. When the window reaches the right-end of the chip, it is re-started at the left-end of the next set of rows, also retaining an overlap with the preceding set of rows.

At each iteration, the algorithm attempts to reduce the total border length inside the window by relocating some of its probes (Fig. 3.4a). First, a random maximal

independent set of spots is selected, and the probes assigned to these spots are removed. The term independent refers to the fact that selected spots can be re-assigned to probes without affecting the border length of other selected spots. The algorithm then creates a bipartite graph with nodes representing the removed probes and the now vacant spots (Fig. 3.4b). The edges of this graph are weighted with the number of border conflicts that are generated by the corresponding assignment. Finally, a minimum weight perfect matching on this graph is computed, and the indicated assignments are made (Fig. 3.4c).

A minimum weight perfect matching requires polynomial time (Gross and Yellen, 2004), but for the small graphs generated by SWM, it can be computed rather quickly. The authors experimentally observed that the best results are obtained with small window sizes (e.g. 6×6) and an overlap of half the window size. More over, employing less effort in each window and executing more cycles of optimization is better than employing more effort in each window and executing less cycles.

Selecting an independent set of spots ensures that the cost of each new assignment can be computed independently of the other assignments. The SWM was designed for border length minimization and it takes advantage of the fact that, in this model, an independent set of spots can be constructed by selecting sites that are not immediate neighbors (spots that do not share a common border). SWM can be adapted for conflict index minimization (to our knowledge, this has not been implemented yet) by using larger windows containing relatively sparse independent sets. Therefore several random independent sets should be constructed before moving the window.

3.5 Row-Epitaxial

Row-Epitaxial (Kahng et al., 2003a) is a variant of the Epitaxial algorithm with two main differences introduced to improve scalability: i) spots are filled in a pre-defined order, namely, from top to bottom, left to right, and ii) only a limited number Q of probes candidates are considered for filling each spot.

Like SWM, Row-Epitaxial improves an initial placement that can be constructed by, for example, Radix-sort + 1-threading. For each spot s with a probe p , it looks at the next Q probes that lie in close proximity (to the right or below s), and swaps p with the probe that generates the minimum number of border conflicts between s and its left and top neighbors.

In the experiments conducted by Kahng et al. (2003a), Row-Epitaxial was the best large-scale placement algorithm (for border length minimization), achieving up to 9% reduction in border conflicts over the TSP + 1-threading, whereas SWM achieved slightly worse results but required significantly less time.

Row-Epitaxial can also be adapted to conflict index minimization by swapping a probe of a spot s with the probe candidate that minimizes the sum of conflict indices in a region around s restricted to those neighboring probes that are to the left or above s (those which have already found their final positions).

Table 3.1 shows the results of using Row-Epitaxial for both border length and conflict index minimization on chips with random probe sequences (uniformly generated). Probes were lexicographically sorted and left-most embedded into the standard 74-step Affymetrix deposition sequence and threaded on the array with k -threading. The resulting layout was then used as a starting point for Row-Epitaxial.

Although Hannenhalli et al. (2002) suggested that 1-threading is the best option for laying out a TSP tour on the chip, our results show that increasing k improves the initial layout produced by sorting + k -threading in terms of border length and conflict index as well. However, the best initial layout does not necessarily lead to the best final layout produced by Row-Epitaxial.

Our results also confirm that the running time of Row-Epitaxial is $O(Qn)$, i.e., linear in the chip size, where Q is a user-defined parameter that controls the number of probe candidates examined for each spot. In this way, solution quality can be traded for running time: More candidates yield better layouts but also demand more time.

3.6 Greedy

As discussed in the previous section, the best results obtained with Row-Epitaxial do necessarily come from the best initial layouts produced by k -threading. This is probably because Row-Epitaxial does not take advantage of the probe order used in k -threading when it looks for probe candidates for filling a certain spot (Row-Epitaxial simply looks in the next Q probes, row-by-row, regardless of how the probes were threaded in the array).

In this section, we present a new placement algorithm, called Greedy, that incorporates the Row-Epitaxial approach with a k -threading filling strategy. Like Row-Epitaxial, spots are filled in a greedy fashion, i.e., for each spot s , the algorithm examines Q probe candidates and chooses the one that can be placed at s with minimum cost (Greedy can also be easily implemented for border length as well as conflict index minimization). There are two main differences, however. First, spots are filled with k -threading instead of simply row-by-row.

Perhaps more importantly, Greedy sorts the probes lexicographically and keeps them in a doubly-linked list, which is used to maintain the probe order during the whole placement. Once a probe p is selected to fill a certain spot, it is removed from the list and the next search of candidates is performed around p 's former position, e.g., $Q/2$

Table 3.1: Normalized border length and average conflict index of layouts produced by Row-Epitaxial (Row-Eptx) on random chips of various dimensions, with initial layouts produced by Radix-sort + k -threading. Running times are reported in minutes and include the time for k -threading and Row-Epitaxial. All results are averages over a set of five chips.

Dim.	Q	k	Border length minimization			Conflict index minimization		
			k -threading	Row-Eptx	Time	k -threading	Row-Eptx	Time
300×300	5K	0	24.9649	18.2935	5.3	701.8698	462.5194	11.8
		1	24.1235	18.2999	5.4	690.8091	462.4656	11.8
		2	23.8695	18.3072	5.3	685.5916	462.6394	11.8
		3	23.7993	18.3226	5.4	683.5980	462.5885	11.8
		4	23.7588	18.3279	5.4	682.3542	462.7775	11.8
	10K	0	24.9649	18.1477	10.5	701.8698	444.0354	22.2
		1	24.1235	18.1529	10.5	690.8091	444.0904	22.4
		2	23.8695	18.1519	10.5	685.5916	444.1960	22.2
		3	23.7993	18.1591	10.5	683.5980	443.9850	22.3
		4	23.7588	18.1603	10.5	682.3542	444.1745	22.2
	20K	0	24.9649	18.0274	20.0	701.8698	426.7824	40.9
		1	24.1235	18.0325	20.0	690.8091	426.8863	40.8
		2	23.8695	18.0277	19.8	685.5916	426.8832	40.8
		3	23.7993	18.0272	20.0	683.5980	426.8694	41.6
		4	23.7588	18.0321	19.9	682.3542	426.6600	40.9
500×500	5K	0	24.2693	17.6000	15.6	693.5428	456.2042	33.5
		1	23.3454	17.6095	15.4	682.2197	456.1167	33.5
		2	23.0797	17.6246	15.5	676.4831	456.4575	33.5
		3	22.9632	17.6474	15.3	672.8211	456.4878	33.7
		4	22.9162	17.6670	15.5	671.2458	456.7849	33.4
	10K	0	24.2693	17.4503	31.4	693.5286	438.6516	64.8
		1	23.3454	17.4523	32.6	682.2197	438.6618	64.7
		2	23.0797	17.4582	32.2	676.4831	438.5759	64.7
		3	22.9632	17.4685	32.0	672.8211	438.7841	65.8
		4	22.9162	17.4755	31.8	671.2458	438.9227	64.6
	20K	0	24.2693	17.3303	63.7	693.5286	421.1773	125.1
		1	23.3454	17.3297	65.5	682.2197	421.0770	125.5
		2	23.0797	17.3308	66.6	676.4831	421.0952	126.0
		3	22.9632	17.3344	66.1	672.8211	420.9112	123.7
		4	22.9162	17.3376	64.8	671.2458	420.9534	124.4
800×800	5K	0	23.6818	16.9760	43.0	689.5915	447.9877	88.2
		1	22.6092	16.9927	43.1	672.2191	448.0745	88.9
		2	22.3205	17.0187	43.3	664.9477	448.6071	89.0
		3	22.1958	17.0589	37.7	660.5476	448.9202	88.3
		4	22.1279	17.1085	37.9	658.4852	449.1434	88.3
	10K	0	23.6818	16.8032	95.0	689.5915	432.2326	179.6
		1	22.6092	16.8111	95.7	672.2191	432.5542	179.3
		2	22.3205	16.8235	95.2	664.9477	432.4991	178.1
		3	22.1958	16.8353	84.6	660.5476	432.6598	179.2
		4	22.1279	16.8622	83.7	658.4852	432.6423	178.9
	20K	0	23.6818	16.6771	219.1	689.5915	415.6571	365.9
		1	22.6092	16.6803	220.1	672.2191	415.7411	367.4
		2	22.3205	16.6851	193.0	664.9477	415.6809	366.9
		3	22.1958	16.6915	191.3	660.5476	415.7138	375.6
		4	22.1279	16.7007	194.0	658.4852	415.7714	368.2

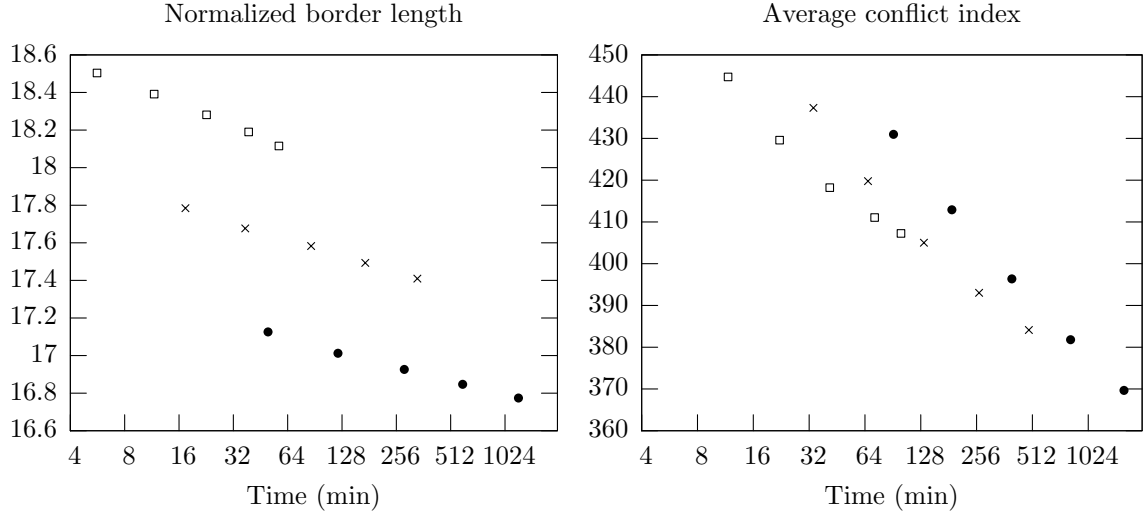


Figure 3.5: Trade-off between solution quality and running time with the Greedy algorithm on random chips of dimensions 300×300 (\square), 500×500 (\times) and 800×800 (\bullet). The number Q of candidates per spot are 5K, 10K, 20K, 40K, and 80K (from left to right). Layouts are measured by normalized border length (left) and average conflict index (right).

probes to the left and to the right of p . In this way, the algorithm always examines the probes that are more likely to be similar to the last placed probe (which will be its neighbor on the chip).

Table 3.2 shows the results of using Greedy for both border length and conflict index minimization on the same set of (random) chips of Table 3.1. Our results show that Greedy is slightly slower than Row-Epitaxial (because it keeps the probe in a doubly-linked list) but marginally better in terms of border length minimization and significantly better in terms of conflict index minimization.

Talk about k -threading.

In terms of border length, we observed that increasing Q above 5K has little positive effect (see Fig. 3.5). For instance, on 800×800 chips, increasing Q from 5K to 10K reduced the normalized border length, on average, by only 0.66% (from 17.1259 to 17.0122). In terms of conflict index, increasing Q even above 40K may result in significant improvements for large chips. For instance, on 800×800 chips, increasing Q from 40K to 80K reduced the average conflict index by 3.18% (from 381.8034 to 369.6583).

TODO: review figures above for border length.

The fact that increasing Q has more effect in terms of conflict index is probably because, in this measure, there is more room for optimization as the conflicts can be moved to the extremities of the probes (while retaining the same number of border conflicts) and a larger number of neighbors are involved.

Table 3.2: Normalized border length (NBL) and average conflict index (ACI) of layouts produced by Greedy on random chips of various dimensions. The results of Row-Epitaxial on the same set of chips (Table 3.1) is shown for comparison. Running times in minutes.

Dim.	Q	k	Border length minimization				Conflict index minimization			
			Row-Epitaxial		Greedy		Row-Epitaxial		Greedy	
			NBL	Time	NBL	Time	ACI	Time	ACI	Time
300 ²	5K	0	24.9649	60.0	18.2935	5.3	701.8698	60.0	462.5194	11.8
		1	24.1235		18.2999	5.4	690.8091		462.4656	11.8
		2	23.8695		18.3072	5.3	685.5916		462.6394	11.8
		3	23.7993		18.3226	5.4	683.5980		462.5885	11.8
		4	23.7588		18.3279	5.4	682.3542		462.7775	11.8
	10K	0	24.9649		18.1477	10.5	701.8698		444.0354	22.2
		1	24.1235		18.1529	10.5	690.8091		444.0904	22.4
		2	23.8695		18.1519	10.5	685.5916		444.1960	22.2
		3	23.7993		18.1591	10.5	683.5980		443.9850	22.3
		4	23.7588		18.1603	10.5	682.3542		444.1745	22.2
	20K	0	24.9649		18.0274	20.0	701.8698		426.7824	40.9
		1	24.1235		18.0325	20.0	690.8091		426.8863	40.8
		2	23.8695		18.0277	19.8	685.5916		426.8832	40.8
		3	23.7993		18.0272	20.0	683.5980		426.8694	41.6
		4	23.7588		18.0321	19.9	682.3542		426.6600	40.9
500 ²	5K	0	24.2693		17.6000	15.6	693.5428		456.2042	33.5
		1	23.3454		17.6095	15.4	682.2197		456.1167	33.5
		2	23.0797		17.6246	15.5	676.4831		456.4575	33.5
		3	22.9632		17.6474	15.3	672.8211		456.4878	33.7
		4	22.9162		17.6670	15.5	671.2458		456.7849	33.4
	10K	0	24.2693		17.4503	31.4	693.5286		438.6516	64.8
		1	23.3454		17.4523	32.6	682.2197		438.6618	64.7
		2	23.0797		17.4582	32.2	676.4831		438.5759	64.7
		3	22.9632		17.4685	32.0	672.8211		438.7841	65.8
		4	22.9162		17.4755	31.8	671.2458		438.9227	64.6
	20K	0	24.2693		17.3303	63.7	693.5286		421.1773	125.1
		1	23.3454		17.3297	65.5	682.2197		421.0770	125.5
		2	23.0797		17.3308	66.6	676.4831		421.0952	126.0
		3	22.9632		17.3344	66.1	672.8211		420.9112	123.7
		4	22.9162		17.3376	64.8	671.2458		420.9534	124.4
800 ²	5K	0	23.6818		16.9760	43.0	689.5915		447.9877	88.2
		1	22.6092		16.9927	43.1	672.2191		448.0745	88.9
		2	22.3205		17.0187	43.3	664.9477		448.6071	89.0
		3	22.1958		17.0589	37.7	660.5476		448.9202	88.3
		4	22.1279		17.1085	37.9	658.4852		449.1434	88.3
	10K	0	23.6818		16.8032	95.0	689.5915		432.2326	179.6
		1	22.6092		16.8111	95.7	672.2191		432.5542	179.3
		2	22.3205		16.8235	95.2	664.9477		432.4991	178.1
		3	22.1958		16.8353	84.6	660.5476		432.6598	179.2
		4	22.1279		16.8622	83.7	658.4852		432.6423	178.9
	20K	0	23.6818		16.6771	219.1	689.5915		415.6571	365.9
		1	22.6092		16.6803	220.1	672.2191		415.7411	367.4
		2	22.3205		16.6851	193.0	664.9477		415.6809	366.9
		3	22.1958		16.6915	191.3	660.5476		415.7138	375.6
		4	22.1279		16.7007	194.0	658.4852		415.7714	368.2

Chapter 4

MLP and the Quadratic Assignment Problem

In this chapter, we show that the Microarray Placement Problem (MLP) with general distance-dependent and position-dependent weights is an instance of the *Quadratic Assignment Problem* (QAP), a classical combinatorial optimization problem introduced by Koopmans and Beckmann (1957), which opens up the way for using QAP techniques to design microarray chips.

We then use an existing QAP heuristic algorithm called GRASP to design the layout of small artificial chips, comparing our results with the best known placement algorithm. The chapter ends with a discussion about how this approach can be combined with other existing algorithms to design and improve larger microarrays.

4.1 Quadratic assignment problem

The quadratic assignment problem (QAP) can be stated as follows. Given $n \times n$ real-valued matrices $F = (f_{ij}) \geq 0$ and $D = (d_{kl}) \geq 0$, find a permutation π of $\{1, 2, \dots, n\}$ such that

$$\sum_{i=1}^n \sum_{j=1}^n f_{ij} \cdot d_{\pi(i)\pi(j)} \rightarrow \min. \quad (4.1)$$

The attribute *quadratic* stems from the fact that the target function can be written with n^2 binary indicator variables $x_{ik} \in \{0, 1\}$, where $x_{ik} := 1$ if and only if $k = \pi(i)$. The objective (4.1) then becomes

$$\sum_{i=1}^n \sum_{j=1}^n f_{ij} \cdot \sum_{k=1}^n \sum_{l=1}^n d_{kl} \cdot x_{ik} \cdot x_{jl} \rightarrow \min,$$

such that $\sum_k x_{ik} = 1$ for all i , $\sum_i x_{ik} = 1$ for all k , and $x_{ik} \in \{0, 1\}$ for all (i, k) . The objective function is a quadratic form in x .

The QAP has been used to model a variety of real-life problems. One common example is the facility location problem where n facilities must be assigned to n locations. The facilities could be, for instance, the clinics, doctors or services (X-ray, emergency room, etc.) provided by a hospital and the locations could be the available rooms of the hospital building.

In this scenario, F is called the *flow matrix* as f_{ij} represents the flow of materials or persons from facility i to facility j . Matrix D is called the *distance matrix*, as d_{kl} gives the distance between locations k and l . One unit of flow is assumed to have an associated cost proportional to the distance between the facilities, and the optimal permutation π defines a one-to-one assignment of facilities to locations with minimum cost.

4.2 QAP formulation of the MLP

The MLP can be seen as an instance of the QAP, where we want to find a one-to-one correspondence between spots and probes minimizing a given penalty function such as total border length or total conflict index (defined in Chapter 2). To formulate it, we use the facility location example by viewing the probes as locations and the spots as facilities, i.e., the spots are assigned to the probes. The flow matrix F then contains the “closeness” values between spots, while the distance matrix D contains the conflicts between probe embeddings.

We first give the general formulation for conflict index minimization case; the border length minimization case is obtained by using the particular weight functions given in Section 2.4.

In a realistic setting, we may have more spots available than probes to place. Below, we show that this does not cause problems as we can add enough “empty” probes and define their weights appropriately.

Perhaps more severely, we assume that all probes have a single pre-defined embedding in order to force a one-to-one relationship. A more elaborate formulation would consider all possible embeddings of a probe, but then it becomes necessary to ensure that only one embedding of each probe is used. This still leads to a quadratic integer programming problem, albeit with slightly different side conditions.

Our goal is to design a microarray minimizing the sum of conflict indices over all spots $s \in \mathcal{S}$, i.e.,

$$\sum_{s \in \mathcal{S}} \mathcal{C}(s) \rightarrow \min .$$

The “flow” f_{ij} between spots i and j depends on their distance on the chip; in accordance with the conflict index model, we set

$$f_{ij} := \mathbb{1}_{\{i,j \text{ neighbors}\}} \cdot \gamma(i, j) \quad (4.2)$$

where “neighbors” means that spots i and j are at most three cells away (horizontally and vertically) from each other. Note that most of the flow values on large arrays are zero. For border length minimization, the case is even simpler: We set $f_{ij} := 1$ if spots i and j are adjacent, and $f_{ij} := 0$ otherwise.

The “distance” d_{kl} between probes k and l depends on the conflicts between their embeddings ε_k and ε_l . To account for possible “empty” probes to fill up surplus spots, we set $d_{kl} := 0$ if k or l or both refer to an empty probe — i.e., empty probes never contribute to the target function since we do not mind if nucleotides are erroneously synthesized on spots assigned to empty probes. For real probes, we set

$$d_{kl} := \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_k, t=0\}} \cdot \omega(\varepsilon_k, t) \cdot \mathbb{1}_{\{\varepsilon_l, t=1\}} \right). \quad (4.3)$$

Note that d_{kl} is related to the conflict index distance $C(k, l)$ defined in Section 2.3 (Eq. 2.10):

$$\begin{aligned} & d_{kl} + d_{lk} \\ &= \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_k, t=0\}} \cdot \omega(\varepsilon_k, t) \cdot \mathbb{1}_{\{\varepsilon_l, t=1\}} \right) + \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_l, t=0\}} \cdot \omega(\varepsilon_l, t) \cdot \mathbb{1}_{\{\varepsilon_k, t=1\}} \right) \\ &= \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_k, t=0 \text{ and } \varepsilon_l, t=1\}} \cdot \omega(\varepsilon_k, t) \right) + \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_l, t=0 \text{ and } \varepsilon_k, t=1\}} \cdot \omega(\varepsilon_l, t) \right) \\ &= \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_k, t=0 \text{ and } \varepsilon_l, t=1\}} \cdot \omega(\varepsilon_k, t) + \mathbb{1}_{\{\varepsilon_l, t=0 \text{ and } \varepsilon_k, t=1\}} \cdot \omega(\varepsilon_l, t) \right) \\ &= C(k, l) \end{aligned}$$

In the case of border length minimization, where $\theta = 0$ and $c = 1/2$ (see Section 2.4), we obtain that $d_{kl} + d_{lk} = H(k, l) = H(l, k)$, where H_{kl} denotes the Hamming distance between the embeddings ε_k and ε_l (Eq. 2.3).

It now follows that for a given assignment π , we have,

$$f_{ij} \cdot d_{\pi(i)\pi(j)} = \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_{\pi(i)}, t=0\}} \cdot \omega(\varepsilon_{\pi(i)}, t) \cdot \mathbb{1}_{\{\varepsilon_{\pi(j)}, t=1\}} \cdot \mathbb{1}_{\{i,j \text{ neighbors}\}} \cdot \gamma(i, j) \right).$$

The objective function (4.1) then becomes

$$\begin{aligned}
 & \sum_i \sum_j f_{ij} \cdot d_{\pi(i)\pi(j)} \\
 &= \sum_i \sum_j \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_{\pi(i),t}=0\}} \cdot \omega(\varepsilon_{\pi(i),t}) \cdot \mathbb{1}_{\{\varepsilon_{\pi(j),t}=1\}} \cdot \mathbb{1}_{\{i,j \text{ neighbors}\}} \cdot \gamma(i,j) \right) \\
 &= \sum_i \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_{\pi(i),t}=0\}} \cdot \omega(\varepsilon_{\pi(i),t}) \cdot \sum_j \mathbb{1}_{\{i,j \text{ neighbors}\}} \cdot \mathbb{1}_{\{\varepsilon_{\pi(j),t}=1\}} \cdot \gamma(i,j) \right) \\
 &= \sum_i \sum_{t=1}^T \left(\mathbb{1}_{\{\varepsilon_{\pi(i),t}=0\}} \cdot \omega(\varepsilon_{\pi(i),t}) \cdot \sum_{j: \text{ neighbor of } i} \mathbb{1}_{\{\varepsilon_{\pi(j),t}=1\}} \cdot \gamma(i,j) \right) \\
 &= \sum_i \mathcal{C}(i),
 \end{aligned}$$

and indeed equals the total conflict index with our definitions of $F = (f_{ij})$ and $D = (d_{kl})$.

Remark. Note that it is technically possible to switch the definitions of F and D , i.e., to assign probes to spots instead of spots to probes as we do now, without modifying the mathematical problem formulation. However, this would lead to high distance values for neighboring spots and many zero distance values for independent spots, a somewhat counterintuitive model. Also, some QAP heuristics initially find pairs of objects with large flow values and place them close to each other. Therefore, the way of modeling F and D may be significant.

4.3 QAP heuristics

We have shown how the microarray placement problem can be modeled as a quadratic assignment problem. However, the QAP is known to be NP-hard and particularly hard to solve in practice. Instances of size larger than $n = 20$ are generally considered to be impossible to solve to optimality. Fortunately, several heuristics exist, including approaches based on tabu search, simulated annealing and genetic algorithms (for a survey, see Çela, 1997). Our formulation is thus of interest because we can now use existing QAP heuristics to design the layout of microarrays minimizing either the sum of border lengths or conflict indices.

As an example, we briefly describe a general QAP heuristic known as GRASP (Li et al., 1994), which was first used for solving the QAP by Feo and Resende (1995),

and an improved version called GRASP with path-relinking (Oliveira et al., 2004), that we used to design small microarray chips with our formulation.

4.3.1 GRASP with Path-relinking

GRASP (Greedy Randomized Adaptive Search Procedure) is comprised of two phases: a construction phase where a random feasible solution is built, and a local search phase where a local optimum in the neighborhood of that solution is sought. In the following description we use the terms of the facility location problem: f_{ij} is the flow between facilities i and j , d_{kl} is the distance between locations k and l .

The construction phase starts by sorting the $(n^2 - n)$ elements of the distance matrix in increasing order and keeping the smallest $E := \lfloor \beta(n^2 - n) \rfloor$ elements, where $0 < \beta < 1$ is a restriction parameter given as input.

$$d_{k_1 l_1} \leq d_{k_2 l_2} \leq \dots \leq d_{k_E l_E}.$$

Similarly, the $(n^2 - n)$ elements of the flow matrix are sorted, this time in decreasing order, and the largest E elements are kept:

$$f_{i_1 j_1} \geq f_{i_2 j_2} \geq \dots \geq f_{i_E j_E}.$$

Then, the costs of assigning pairs of facilities to pairs of locations are computed. The cost of initially assigning facility i_q to location k_q and facility j_q to location l_q for some $q \in \{1, \dots, E\}$ is $d_{k_q l_q} f_{i_q j_q}$. GRASP sorts the vector

$$(d_{k_1 l_1} f_{i_1 j_1}, d_{k_2 l_2} f_{i_2 j_2}, \dots, d_{k_E l_E} f_{i_E j_E}),$$

keeping the $\lfloor \alpha E \rfloor$ smallest elements, where $0 < \alpha < 1$ is another restriction parameter. A simultaneous assignment of a pair of facilities to a pair of locations is selected at random among those with the $\lfloor \alpha E \rfloor$ smallest costs, and a feasible solution is then built by making a series of greedy assignments.

In the local search phase, GRASP searches for a local optimum in the neighborhood of the constructed solution. Several search strategies and definitions of neighborhood can be used. One possible approach is to check every possible swap of assignments and make only those which improve the current solution until no further improvements can be made.

The construction and local search phases are repeated for a given number of times, and the best solution found is returned.

Path-relinking. GRASP takes no advantage of the knowledge gained in previous iterations to build or improve an obtained solution, i.e., each new solution is built from scratch.

GRASP with path-relinking is an extension of the basic GRASP algorithm that uses an “elite set” to store the best solutions found. It incorporates a third phase that chooses, at random, one elite solution that is used to improve the solution produced at the end of the local search phase.

Solutions p and q are combined as follows. For every location $k = 1, \dots, n$, the path-relinking algorithm attempts to exchange facility p_k assigned to location k in solution p with facility q_k assigned to location k in the elite solution. In order to keep the solution p feasible, it exchanges p_k with p_l , where $p_l = q_k$. This exchange is performed only if it results in a better solution. The result of the path-relinking phase is a solution r that is at least as good as the better of p and q .

4.4 Results

We present experimental results of using GRASP with path-relinking (GRASP-PR) for designing the layout of small artificial chips, and compare them with the layouts produced by Row-Epitaxial (described in Chapter 3).

We used a C implementation of GRASP-PR provided by Oliveira et al. (2004) with default parameters (32 iterations, $\alpha = 0.1$, $\beta = 0.5$, and elite set of size 10). The main routine takes three arguments: the dimension n of the problem (in our case, the number of spots or probes) and matrices F and D . The matrices were generated using the formulations presented in Section 4.2.

The data set consists of chips with probes of length 25 uniformly generated and asynchronously embedded in a deposition sequence of length 74. The running times and the border lengths of the resulting layouts are shown in Table 4.1 (all results are averages over a set of ten chips).

Our results show that GRASP-PR produces layouts with lower border lengths than Row-Epitaxial on the smaller chips. On 6x6 chips, GRASP-PR outperforms Row-Epitaxial by 2.14 percentage points on average (15.94% – 13.80%), when compared to the initial random layout. On 9x9 chips, however, this difference drops to 0.16 percentage points, while Row-Epitaxial generates better layouts on 11x11 or larger chips. In terms of running time, Row-Epitaxial is faster and shows little variation as the number of probes grows. In contrast, the time required to compute a layout with GRASP-PR increases at a fast rate.

Table 4.2 shows better results in terms of conflict indices. GRASP-PR consistently produces better layouts on all chip dimensions, achieving up to 6.38% less conflicts

Table 4.1: Border length of random chips compared with the layouts produced by Row-Epitaxial and GRASP with path-relinking. Reductions in border length are reported in percentages compared to the random layout.

Chip dimension	Random	Row-Epitaxial			GRASP with path-relinking		
	Border length	Border length	Reduction (%)	Time (sec.)	Border length	Reduction (%)	Time (sec.)
6 × 6	1 989.20	1 714.60	13.80	0.01	1 672.20	15.94	2.73
7 × 7	2 783.20	2 354.60	15.40	0.02	2 332.60	16.19	6.43
8 × 8	3 721.20	3 123.80	16.05	0.03	3 099.13	16.72	12.49
9 × 9	4 762.00	3 974.80	16.53	0.05	3 967.20	16.69	25.96
10 × 10	5 985.20	4 895.60	18.20	0.06	4 911.40	17.94	47.57
11 × 11	7 288.40	5 954.40	18.30	0.10	5 990.73	17.80	87.48
12 × 12	8 714.00	7 086.20	18.68	0.11	7 159.80	17.84	152.42

Table 4.2: Average conflict indices of random chips compared with the layouts produced by Row-Epitaxial and GRASP with path-relinking.

Chip dimension	Random	Row-Epitaxial			GRASP with path-relinking		
	Avg. C. Index	Avg. C. Index	Reduction (%)	Time (sec.)	Avg. C. Index	Reduction (%)	Time (sec.)
6 × 6	524.28	495.15	5.56	0.05	467.08	10.91	3.68
7 × 7	558.25	521.90	6.51	0.07	489.32	12.35	8.84
8 × 8	590.51	551.84	6.55	0.09	515.69	12.67	19.48
9 × 9	613.25	568.62	7.28	0.11	533.79	12.96	38.83
10 × 10	628.50	576.49	8.28	0.11	539.69	14.13	73.09
11 × 11	642.72	588.91	8.37	0.12	551.41	14.21	145.67
12 × 12	656.86	598.21	8.93	0.12	561.21	14.56	249.19

on 10x10 chips, for example, when compared to Row-Epitaxial. In terms of running times, GRASP-PR is even slower than in the border length case. The reason is not clear, but it could be related to the fact that the distance matrix contains fewer zero entries with the conflict index formulation.

The gains in terms of conflict index of both Row-Epitaxial and GRASP-PR are clearly less than the gains in terms of border length (when compared to the initial random layout). This may be because the embeddings are fixed and the reduction of conflicts is restricted to the relocation of the probes, which only accounts for one part of the conflict index model.

4.5 Discussion

The QAP is notoriously hard to solve, and currently known exact methods start to take prohibitively long already for slightly more than 20 objects, i.e., we could barely solve the problem exactly for 5×5 arrays. Fortunately, the literature on QAP heuristics is rich, as many problems in operations research can be modeled as QAPs. Here we used one such heuristic to identify the potential of the MLP-QAP-relation.

As our results show, however, even heuristic algorithms are too slow to deal with chips of dimensions larger than 12×12 , and although we could design a 20×20 chip with a QAP heuristic within a day, we have to keep in mind that this would still be a very small part of bigger problem as real microarray dimensions range from 200×200 up to 1164×1164 .

For this reason, we restricted our experiments to such small chips and QAP heuristics that could handle the problems within a few minutes. Up to now, finding exact solutions even to these small microarrays seems to be an incredible hard task. We mention here experiments conducted by Dr. Peter Hahn, who used two branch-and-bound algorithms to solve some problem instances from Table 4.1. With RTL-2 (Adams et al., To appear), it was possible to find two solutions with total border length of 1 652 for a selected 6×6 chip, being only 1.43% better than the solution found with GRASP-PR (1 676), although it took RTL-2 about 6.5 hours, in contrast with the less than 3 seconds needed by GRASP-PR. A lower bound calculation for the same problem resulted in 1 624, so the RTL-2 solution is only 1.69% higher, while the gap to the GRASP-PR solution is about 3.10%.

For another selected problem of dimension 7×7 , Dr. Hahn found one solution with border length 2 290 using RTL-1 (Hahn et al., 1998), being about 1.72% better than the solution found by GRASP-PR (2 330), although it took RTL-1 some 29 hours, in contrast with the less than 7 seconds needed for the GRASP-PR run. The results obtained with exact QAP solvers give an idea of how hard the quadratic assignment problem actually is, and show that the results with GRASP-PR are a good compromise when time is limited.

Improved results for several selected problem instances from Tables 4.1 and 4.2 were also reported by Chris MacPhee using GATS, a hybrid genetic / tabu search algorithm, although these results were obtained on a number of large memory SMP machines, each having 144 processors and 576 GB of global memory. The latest results for these selected problems are available online at <http://gi.cebitec.uni-bielefeld.de/assb/chiplayout/qap>.

4.5.1 Alternatives

It is clear that, because of the large number of probes on industrial microarrays, it is not feasible to use GRASP-PR (or any other currently available QAP method) to design an entire microarray chip. However, we showed that it is certainly possible to use it on small sub-regions of a chip, which opens up the way for two alternatives.

First, the QAP approach could be used combined with a partitioning algorithm such as those discussed in Chapter 6 to design the smaller regions that result from the partitioning. This, however, does not seem promising because, as we will see later, a partitioning is a compromise in solution quality, and level of partitioning required to achieve the dimensions supported by the QAP approach is too high.

It is interesting to extrapolate the times shown on Table 4.1 to predict the total time that would be required to design the layout of commercial microarrays, if we were to combine GRASP-PR with a partitioning algorithm. If the partitioning produced 6×6 regions, 37 636 sub-regions would be created from the 1164×1164 Affymetrix Human Genome U133 Plus 2.0 GeneChip array, one of the largest Affymetrix chips. Since each sub-region takes around 3 seconds to compute with GRASP-PR, the total time required for designing such a chip would be a little over 31 hours (ignoring the time for the partitioning itself).

If the partitioning produced 12×12 regions, 9 409 sub-regions would be created and, at 2.4 minutes each, the total time would be more than 16 days. This is probably prohibitive, although it is certainly possible to reduce the time of each GRASP-PR execution by running it on faster machines or run them in parallel.

A better alternative is to use the QAP approach to improve an existing layout, iteratively, by relocating probes inside a defined region of the chip, in a sliding-window fashion. Each iteration of this method would produce an instance of a QAP whose size equals the number of spots inside the window. The QAP heuristics could then be used to check whether a different arrangement of the probes inside the window can reduce the conflicts. For this approach to work, however, we also need to take into account the conflicts due to the spots around the window. Otherwise, a new layout with less internal conflicts could be achieved at the expense of increasing conflicts on the borders of the window.

A simple way of preventing this problem is to solve a larger QAP instance consisting of the spots inside the window as well as those in a layer (of three spots) around it. The spots outside the window obviously must remain unchanged, and that can be done by fixing the corresponding elements of the permutation π . Note that there is no need to compute f_{ij} if spots i and j are both outside the window, nor d_{kl} if probes k and l are assigned to spots outside the window.

Chapter 5

Re-embedding Algorithms

Once the probes have been placed, conflicts can be further reduced by re-embedding the probes without changing their locations. All re-embedding algorithms presented in this section are based on the Optimum Single Probe Embedding (OSPE) introduced by Kahng et al. (2002). OSPE is a dynamic programming for computing an optimum embedding of a single probe with respect to its neighbors, whose embeddings are considered as fixed. The algorithm was originally developed for border length minimization but here we present a more general form designed for the conflict index model (de Carvalho Jr. and Rahmann, 2006a).

5.1 Optimum Single Probe Embedding

The OSPE algorithm can be seen as a special case of a global alignment between a probe sequence p of length ℓ and the deposition sequence N of length T , disallowing mismatches and gaps in N . We assume that p is placed at spot s , and that we know the embeddings of all probes placed at spots near s .

The optimal embedding of p into N is built by determining the minimum cost of embedding a prefix of p into a prefix of N : We use an $(\ell + 1) \times (T + 1)$ matrix D , where $D[i, t]$ is defined as the minimum cost of an embedding of $p[1..i]$ into $N[1..t]$. The cost is the sum of conflicts induced by the embedding of $p[1..i]$ on its neighbors, plus the conflicts suffered by $p[1..i]$ because of the embeddings of its neighbors.

We can compute the value for $D[i, t]$ by looking at two previous entries in the matrix: $D[i, t - 1]$ and $D[i - 1, t - 1]$. The reason is that $D[i, t]$ is the minimum cost of embedding $p[1..i]$ up to the t -th synthesis step, which can only be obtained from the previous step ($t - 1$) by either masking or unmasking spot s at step t .

If s is productive at step t , base N_t is appended to $p[1..i - 1]$; this is only possible if $p[i] = N[t]$. In this case a cost U_t is added for the risk of damaging probes at neighboring spots s' . We know that $p[1..i - 1]$ can be embedded in $N[1..t - 1]$ with

optimal cost $D[i - 1, t - 1]$. Hence, the minimum cost at step t , if s is productive, is $D[i - 1, t - 1] + U_t$. According to the conflict index model,

$$U_t := \sum_{\substack{s': \text{ neighbor} \\ \text{of } s}} \mathbb{1}_{\{\varepsilon_{k(s'), t} = 0\}} \cdot \omega(\varepsilon_{k(s'), t}) \cdot \gamma(s', s).$$

If s is masked at step t , no base is appended to $p[1..i]$, but a cost $M_{i,t}$ must be added for the risk of damaging p (by light directed at neighboring spots s'). Since $D[i, t - 1]$ is the minimum cost of embedding $p[1..i]$ in $N[1..t - 1]$, the minimum cost up to step t , if s is unmasked, is $D[i, t - 1] + M_{i,t}$.

Note that $M_{i,t}$ depends on the number of bases p already contains (that is, on i): Each unmasked neighboring spot s' generates a conflict on p with cost $\gamma(s, s') \cdot c \cdot \exp[\theta \cdot (1 + \min\{i, \ell - i\})]$, in accordance with (2.6)–(2.8). Thus

$$M_{i,t} := c \cdot \exp[\theta \cdot (1 + \min\{i, \ell - i\})] \cdot \sum_{\substack{s': \text{ neighbor} \\ \text{of } s}} \mathbb{1}_{\{\varepsilon_{k(s'), t} = 1\}} \cdot \gamma(s, s').$$

Finally, $D[i, t]$ is computed as the minimum cost of the possible actions,

$$D[i, t] := \begin{cases} \min\{D[i, t - 1] + M_{i,t}, D[i - 1, t - 1] + U_t\} & \text{if } p[i] = N[t], \\ D[i, t - 1] + M_{i,t} & \text{if } p[i] \neq N[t]. \end{cases}$$

The first column of D is initialized as follows: $D[0, 0] = 0$ and $D[i, 0] = \infty$ for $0 < i \leq \ell$, since no probe of length $\ell > 0$ can be embedded into an empty deposition sequence. The first row is initialized by setting $D[0, t] = D[0, t - 1] + M_{0,t}$ for $0 < t \leq T$.

If we assume that costs U_t and $M_{i,t}$ can be computed in constant time, the time complexity of the OSPE algorithm is $O(\ell T)$ since there are $O(\ell T)$ entries in D to compute. The algorithm can be rather time-consuming in the general form presented here, since we have to look at the embeddings of up to 48 neighbors around s . Naturally, it runs much faster for border length minimization, since there are only four neighbors, and there are neither position-dependent (ω) nor distance-dependent (γ) weights to compute. In any case, a few optimizations significantly reduce the running time. For instance, in each row, only the columns between the left-most and the right-most embedding of p in N need to be computed.

Once D is computed, the minimum cost is $D[\ell, T]$, and an optimal embedding of p into N can be constructed by tracing a path from $D[\ell, T]$ back to $D[0, 0]$, similarly to the procedure used to build an optimal global alignment. This takes $O(T)$ time.

5.2 Re-embedding algorithms

The OSPE algorithm is the basic operation of several post-placement optimization algorithms: Greedy, Batched Greedy and Chessboard (Kahng et al., 2002), and Sequential (Kahng et al., 2003b). Their main difference lies in the order in which the probes are re-embedded.

Since OSPE never increases the amount of conflicts in the region around the re-embedded probe, optimization algorithms can execute several re-embedding operations without risk of worsening the current solution. Moreover, probes can be re-embedded several times since new improvements may be possible once neighbors are changed. In fact, the following algorithms work in repeating cycles of optimization until no more improvements are possible (when a local optimal solution is found), or until improvements drop below a given threshold.

The Greedy algorithm uses OSPE to compute, for each spot of the chip, the maximum reduction of border conflicts achievable by optimally re-embedding its probe. It then selects a spot s with the highest gain (reduction of conflicts) and re-embeds its probe optimally, updating the gains of affected neighboring spots.

A faster version of this algorithm, called Batched Greedy, pre-selects several spots for re-embedding and thus sacrifices its greedy nature by postponing the update of gains.

The Chessboard optimization is based on the fact that a chip can be bi-colored like a chessboard, in such a way that the embeddings of probes located on white spots are independent of those placed on black spots (with respect to border length), and vice-versa. The Chessboard uses this coloring to alternate the optimal re-embedding of probes located on black and white spots.

The Sequential optimization is the simplest algorithm among the four. It proceeds spot by spot, from top to bottom, from left to right, re-embedding each probe optimally. Once the end of the array is reached, it restarts at the top left for the next iteration.

Surprisingly, the Sequential algorithm achieves the greatest reduction of border conflicts with a running time comparable to Batched Greedy, the fastest among the four. All re-embedding algorithms mentioned here were initially developed for border length minimization, but they can all be applied to the conflict index model as well. For the Chessboard optimization, $4 \times 4 = 16$ colors must be used instead of 2.

Chapter 6

Partitioning Algorithms

We mentioned earlier that the MLP is usually approached in two phases: placement and re-embedding. The placement, however, is sometimes preceded by a *partitioning* phase, which breaks the problem into smaller sub-problems that are easier to manage. This is especially helpful for placement algorithms with non-linear time or space complexities that are otherwise unable to handle very large chips.

A partitioning algorithm divides the set of probes \mathcal{P} into smaller subsets, and assigns them to defined regions of the chip. Each region can then be treated as an independent chip (and processed by a placement algorithm) or recursively partitioned. Linear-time placement algorithms may also benefit from a partitioning since probes with similar embeddings are typically assigned to the same region (Row-Epitaxial, for instance, is more likely to find good candidates for filling a spot).

We describe four partitioning algorithms: 1-Dimensional Partitioning, 2-Dimensional Partitioning, Centroid-based Quadrisection (CQ), and Pivot Partitioning (PP). Like placement algorithms, they assume that an initial (left-most, right-most, synchronous or otherwise pre-computed) embedding of the probes is given. Pivot Partitioning is the only algorithm that modifies these embeddings. As we shall see, 1-D and 2-D Partitioning generate a few masks with extremely few conflicts, leaving the remaining masks with high levels of conflicts that are difficult to handle. CQ and PP offer a more uniform optimization over all masks. Results of de Carvalho Jr. and Rahmann (2006a) indicate that PP produces better layouts than CQ on large chips.

Partitioning is a compromise in solution quality since it restricts the space of solutions and may lead to conflicts at partition borders. However, it can improve solution quality in practice when the placement algorithm cannot handle large regions well. It is not advisable to perform too many levels of partitionings because smaller sub-regions mean less freedom for optimization during placement. The right balance depends on both the placement algorithm and the partitioning algorithm.




Figure 6.1: First four levels of 1-Dimensional Partitioning. Dashed lines show the divisions performed in each step; solid lines indicate regions delimited in previous steps (there are no border conflicts between spots separated by solid lines). Masked (shaded) regions are labeled “0”, unmasked (white) regions are labeled “1”. This labeling forms a binary Gray code (shown in the first three steps only).

6.1 1-Dimensional Partitioning

1-Dimensional Partitioning divides the set of probes based on the state of their embeddings at a particular synthesis step. It starts by creating two subsets of \mathcal{P} :

$$\mathcal{P}_0 = \{p_k \in \mathcal{P} | \varepsilon_{k,1} = 0\}, \quad \mathcal{P}_1 = \{p_k \in \mathcal{P} | \varepsilon_{k,1} = 1\}.$$

In other words, \mathcal{P}_0 contains all probes whose embeddings are unproductive during the first synthesis step, whereas \mathcal{P}_1 contains the probes with productive embeddings. The chip is then divided into two horizontal bands, proportionally to the number of probes in \mathcal{P}_0 and \mathcal{P}_1 , so each band accommodates one subset of \mathcal{P} .

This procedure is recursively applied to each band, using the the next synthesis steps to further divide each subset of probes. For instance, the following subsets of \mathcal{P}_0 and \mathcal{P}_1 are created during step $t = 2$:

$$\begin{aligned} \mathcal{P}_{00} &= \{p_k \in \mathcal{P}_0 | \varepsilon_{k,2} = 0\}, & \mathcal{P}_{01} &= \{p_k \in \mathcal{P}_0 | \varepsilon_{k,2} = 1\}, \\ \mathcal{P}_{10} &= \{p_k \in \mathcal{P}_1 | \varepsilon_{k,2} = 0\}, & \mathcal{P}_{11} &= \{p_k \in \mathcal{P}_1 | \varepsilon_{k,2} = 1\}. \end{aligned}$$

The next assignments of subsets to the upper or lower band of their regions are made in such a way that regions with the same “state” — productive (unmasked) or unproductive (masked) — are joined as far as possible, resulting in masks that consist of alternating layers of masked and unmasked spots. This process is illustrated in Fig. 6.1, where at each step t , a band is labeled “0” when its embeddings are unproductive, and “1” when its embeddings are productive. The resulting binary numbers from top to bottom form a binary Gray code, which is defined as a permutation of the binary numbers between 0 and $2^n - 1$ such that neighboring elements have exactly one differing bit, as do the first and last elements (Kreher and Stinson, 1999).

The Gray code highlights an interesting property of 1-D Partitioning. After d levels of partitionings (based on steps 1 to d), the embeddings of any two immediate neighbors differ among the first d steps in at most one step. As a result, masks $M_1 \dots M_d$ exhibit a layered structure that effectively reduces border conflicts.

Unfortunately, the Gray code is disrupted as soon as a region cannot be divided (because all probes of that region are, for instance, masked at a particular step). This will certainly happen as several binary numbers are unlikely to be substrings of embeddings (think of, for example, a long run of zeros).

Moreover, 1-D Partitioning can optimize only a limited number of masks because the sub-regions soon become too narrow to be further divided. The maximum *partitioning depth* d_{max} is primarily limited by the number of rows in the chip. In practice, since regions are likely to be unevenly divided, d_{max} varies between regions. The algorithm can also be configured to stop partitioning a region once its dimensions drop below a given threshold.

1-D Partitioning is easier to implement if the partitionings always produce rectangular regions (i.e., splitting a row between two regions is not allowed). In order to force an exact division of a region, however, it might be necessary to move a few probes from one subset of probes to the other one.

For example, imagine that a chip with $|\mathcal{P}| = 900$ probes, $n_r = 30$ rows and $n_c = 30$ columns is to be partitioned based on the state of the embeddings at the first synthesis step, resulting in sub-sets \mathcal{P}_0 and \mathcal{P}_1 with, say, 638 and 262 probes, respectively. The chip must thus be divided into two sub-regions, the upper one containing $\lceil 30 \cdot 638/900 \rceil = 21$ rows and the lower one with $\lceil 30 \cdot 262/900 \rceil = 9$ rows ($\lceil x \rceil$ is x rounded to the nearest integer). The problem is that the upper region then contains $21 \cdot 30 = 630$ spots but it has to accommodate 638 probes, whereas the lower region contains $9 \cdot 30 = 270$ spots but only 262 probes. The solution is to (arbitrarily) move 8 probes from \mathcal{P}_0 to \mathcal{P}_1 , which, results in some imperfections in the layers of the corresponding mask (a few masked spots in a region of unmasked spots, for instance).

6.2 2-Dimensional Partitioning

The 2-Dimensional Partitioning algorithm extends the idea of 1-D Partitioning to two dimensions, with the potential of optimizing twice as many masks. The algorithm is similar: \mathcal{P} is divided into subsets based on the state of the embeddings at a particular synthesis step. The difference is that 2-D Partitioning alternates horizontal and vertical divisions of regions, and that the assignments of probes to regions obey a two-dimensional binary Gray code (Fig. 6.2).




Figure 6.2: First four levels of 2-Dimensional Partitioning. Dashed lines show the divisions performed in each step; solid lines indicate regions delimited in previous steps. Masked regions are labeled with “0”, unmasked regions with “1”; this labeling forms an approximation to a two-dimensional binary Gray code.

In a 2-D Gray code, two neighboring numbers differ in at most one bit. Thus, regions whose embeddings are at the same state (productive or unproductive) are joined as far as possible.

If regions were always equally divided, 2-D Partitioning would have the same property as 1-D Partitioning: After d levels of partitionings (based on steps 1 to d), the embeddings of any two immediate neighbors would differ among the first d steps in at most one step. However, this is not always the case since 2-D Partitioning is likely to create regions with different dimensions, forcing some regions to share a border with more than its four natural neighbors (for example, in Fig. 6.2 region “0010” borders with “0001” and “1011”).

So far we have described both 1-D and 2-D Partitionings using the state of the first d synthesis steps to divide the set of probes. The result of this approach is that, while the first masks are optimized, the remaining masks are left with high levels of border conflicts; we call this a *left-most mask optimization*.

However, a defect in the middle of the probe is more harmful than in its extremities, so it is more important to optimize the central masks, which synthesize the middle bases. Thus we partition the chip based on the following sequence of synthesis steps, assuming that T is even and d is odd: $T/2, (T/2) \pm 1, (T/2) \pm 2, \dots, (T/2) \pm \lfloor d/2 \rfloor$; we call this a *centered mask optimization*.

For left-most optimization, it makes sense to embed the probes in a left-most fashion in order to reduce conflicts in the last masks (which are not optimized by the partitioning); the left-most embeddings reduce the number of unmasked spots in the last steps, resulting in masks that largely consist of masked spots. Similarly, centered mask optimization produces better results with *centered embeddings*. A centered embedding is constructed by shifting a left-most embedding to the right so that the number of masked steps to the left of the first productive step approximately equals the number of masked steps to the right of the last productive step.




Figure 6.3: Selected masks generated by 2-D Partitioning for a random 300×300 chip with 25-mer probes leftmost embedded into the standard Affymetrix deposition sequence. Unmasked spots are represented by red dots, masked spots by black dots.




Figure 6.4: Normalized border length (on the y-axis) per masking step (on the x-axis) of a layout produced by 2-Dimensional Partitioning for a 1000×1000 chip with random probe sequences (embedded in the standard 74-step Affymetrix deposition sequence). Partitioning stops when a region becomes smaller than 64×64 ; Row-Epitaxial is used for the placement. (x) Left-most mask optimization with left-most embeddings; (□) centered mask optimization with centered embeddings.

Fig. 6.4 shows the results of 2-D Partitioning on a 1000×1000 chip with both optimizations. For left-most mask optimization, we obtain a normalized border length of 33.89 (up to approximately 0.6 per step). For centered mask optimization, the normalized border length improves slightly to 33.59. The average conflict index (not shown in the figure) for left-most mask optimization is 571.8; for centered mask optimization, it improves considerably to 383.5 because of the higher weight of the middle bases.

6.3 Centroid-based Quadrisection

Centroid-based Quadrisection or CQ (Kahng et al., 2003b) employs a different criterion for dividing the set of probes and a different approach for partitioning. At each iteration, a region R is quadrisected into R_1 , R_2 , R_3 , and R_4 . Each sub-region R_i is associated with a selected probe $p_{c_i} \in \mathcal{P}$, called *centroid*, that is used to guide the assignment of the remaining probes to the sub-regions.

A centroid is a representative of its region; it should symbolize the “average embedding” in that region. The remaining probes $p_k \in \mathcal{P} \setminus \{p_{c_1}, p_{c_2}, p_{c_3}, p_{c_4}\}$ are compared to each centroid and assigned to the sub-region R_i whose centroid’s embedding ε_{c_i} has minimum $H(k, c_i)$, where $H(k, k')$ is the Hamming distance between the embeddings ε_k of p_k and $\varepsilon_{k'}$ of $p_{k'}$ (i.e., the number of steps in which ε_k and $\varepsilon_{k'}$ differ).

In order to improve the clustering of similar probes, the four centroids should be very different from each other. The following heuristic is used: First, a probe index c_1 is randomly selected from $\{1, \dots, |\mathcal{P}|\}$. Then, a probe index $c_2 \neq c_1$ maximizing




Figure 6.5: First three levels of Centroid-base Quadrisecion Partitioning. Dashed lines show the divisions performed in each step; solid lines indicate regions delimited in previous steps. The centroids of each partition $R_1 \dots R_4$ are represented by small circles (labeled with $q_1 \dots q_4$ in the first step).

$H(c_2, c_1)$ is selected. Similarly, c_3 maximizing $H(c_3, c_1) + H(c_3, c_2)$ and c_4 maximizing $H(c_4, c_1) + H(c_4, c_2) + H(c_4, c_3)$ are selected. The assignment of centroids to the quadrisecions of the chip is arbitrary.

In order to recover from a possibly bad choice of centroids, one can use a “multi-start heuristic”, running the centroid selection procedure several times (using different “seeds” for c_1), and keeping those that lead to the best partitioning (partitioning quality is measured by the sum of Hamming distances of probe embeddings to their corresponding centroid embeddings).

The partitioning continues recursively until a pre-defined depth has been reached.

CQ was developed for border length minimization, but can be adapted for conflict index minimization by using the *conflict index distance* $C(k, k')$ instead of the Hamming distance $H(k, k')$ between the embeddings ε_k and $\varepsilon_{k'}$.

6.4 Pivot Partitioning: Merging partitioning and re-embedding

Pivot Partitioning or PP (de Carvalho Jr. and Rahmann, 2006a) is, to a certain extent, similar to CQ: Sub-regions are recursively associated with special probes p_{c_i} , here called *pivots* instead of centroids, that are used to guide the assignment of the other probes to the sub-regions. The main differences between PP and CQ are as follows.

Instead of quadrisecioning the chip, PP creates sub-regions by alternating horizontal and vertical divisions (like 2-D Partitioning). The advantage is that regions are divided

proportionally to the size of each subset of probes, so they are not required to have the same size. Furthermore, for each partitioning level, only two pivots need to be selected.

Another distinction is motivated by the fact that different probes have different numbers of embeddings, ranging from a single one to several millions. Probes with more embeddings can more easily adapt to the other probes, that is, they are more likely to have an embedding with fewer conflicts to fill a particular spot than a probe that has only a limited number of embeddings. For this reason, PP uses probes with a single embedding (or few embeddings) as pivots, and chooses the other probes' embeddings and region assignments accordingly.

Indeed, the most important feature of PP is the simultaneous embedding and assignment of probes to sub-regions. Let $M(k, c_i)$ denote the minimum conflict index distance $C(k, c_i)$, as defined in (2.10), over all embeddings of p_k ; we call it the *minimum conflict index distance* between probes p_k and p_{c_i} . It can be efficiently computed with a variant of the OSPE algorithm that ignores the location of the probes and the distance-dependent weights γ . Now, a non-pivot probe p_k is assigned to the region R_i whose pivot p_{c_i} has minimum $M(k, q_i)$ over $i = 1, 2$. Pivot Partitioning continues recursively up to a pre-defined depth. Finally, each probe is embedded to minimize conflicts with its assigned pivot.

Chapter 7

Merging Placement and Re-embedding

The problem with the traditional “place and re-embed” approach is that the arrangement of probes on the chip is decided based on embeddings that are likely to change during the re-embedding phase. Intuitively, better results should be obtained when the placement and embedding phases are considered simultaneously instead of separately. However, because of the generally high number of embeddings of each single probe, it is not easy to design algorithms that efficiently use the additional freedom and run reasonably fast in practice.

We describe Greedy+, the first placement algorithm that simultaneously places and re-embeds the probes, and compare it with Row-Epitaxial, the best known large-scale placement algorithm.

7.1 Greedy+

The goal is to design an algorithm that is similar to Row-Epitaxial, so that we can make a better assessment of the gains resulting from merging the placement and re-embedding phases.

Greedy+ fills the spots row-by-row, from left to right, in a greedy fashion, similarly to Row-Epitaxial. Also, for each spot s , it looks at Q probe candidates and chooses the one that can be placed at s with minimum cost. The difference is that we now consider all possible embeddings of a candidate p instead of only p ’s initial embedding. This is done by temporarily placing p at s and computing its optimal embedding with respect to the already-filled neighbors of s (using OSPE from Sec. 5.1).

Compared to Row-Epitaxial, Greedy+ spends more time evaluating each probe candidate p for a spot s . While Row-Epitaxial takes $O(T)$ time to compute the conflict index or the border length resulting from placing p at s , Greedy+ requires $O(\ell T)$ time since it uses OSPE (recall that ℓ is the probe length and T is the length of the

deposition sequence). To achieve a running time comparable to Row-Epitaxial, we must therefore consider lower candidate numbers Q .

There are a few optimizations that reduce the time spent with OSPE computations when several probe candidates are examined in succession for the same spot. First, we note that if two probe candidates p and p' share a common prefix of length l , the first $l + 1$ rows of the OSPE dynamic programming matrix D will be identical. In other words, if we have calculated the minimum cost of p , we can speed up the calculation of the minimum cost of p' by skipping the first $l + 1$ rows of D .

In order to fully exploit this fact, we examine the probes in lexicographical order so that we maximize the length of the common prefix between two consecutive candidates. We keep a doubly-linked list of probes and remove a probe p from the list when it is placed. For the next spot to be filled, we look at Q probes in the list around p 's former position, e.g., at $Q/2$ probes to the left and to the right of p .

Second, the U_t costs of OSPE need to be computed only once for a given spot s since U_t does not depend on the probe placed at s . Thus, in order to examine another candidate, we only need to recompute the $M_{i,t}$ costs.

Finally, once we know that a probe candidate p can be placed at s with minimum cost C , we can stop the OSPE computation for another candidate p' as soon as all values in a row of D are greater than or equal to C .

7.2 Results

We compare the results of Greedy+ with Row-Epitaxial. To be fair, since Row-Epitaxial is a traditional placement algorithm that does not change the embeddings, we need to compare the layouts obtained by both algorithms after a re-embedding phase. For this task we use the Sequential algorithm (Sec. 5.2) with thresholds of $W = 0.1\%$ for border length minimization and $W = 0.5\%$ for conflict index minimization, so that the algorithm stops as soon as the improvement in one iteration drops below W .

Table 7.1 shows the total border length and the average conflict index of layouts produced by both algorithms on two chips with dimensions 335×335 and 515×515 , filled with probes randomly selected from existing GeneChip arrays (E.Coli Genome 2.0 and Maize Genome, respectively). Probes are initially left-most embedded into the standard 74-step Affymetrix deposition sequence $\{\text{TGCA}\}^{18}\text{TG}$. The parameter Q is chosen differently for both algorithms so that the running time is approximately comparable (e.g., for border length minimization, $Q = 350$ for Greedy+ corresponds to $Q = 10\,000$ for Row-Epitaxial). We make the following observations.

First, increasing Q linearly increases placement time, while only marginally improving chip quality for border length minimization.

Table 7.1: Normalized border length (NBL) and average conflict index (ACI) of layouts produced by Row-Epitaxial and Greedy+ placement (Pl.), followed by Sequential re-embedding (Re-emb.) with thresholds $W = 0.1\%$ for border length minimization and $W = 0.5\%$ for conflict index minimization. Q is the number of probe candidates considered for each spot during placement. Running times are given in seconds.

Border length min.	335 × 335 (E.Coli)		515 × 515 (Maize)	
Row-Epitaxial. Q	10 K	20 K	10 K	20 K
Time (Pl. + Re-emb.)	629 + 9	1211 + 9	3333 + 38	6806 + 38
NBL (Pl.)	34.11	33.81	33.11	32.87
NBL (Pl. + Re-emb.)	33.93	33.66	32.95	32.73
Greedy+. Q	350	700	350	700
Time (Pl. + Re-emb.)	596 + 12	1158 + 12	2633 + 53	4974 + 53
NBL (Pl.)	33.79	33.22	33.07	32.38
NBL (Pl. + Re-emb.)	33.53	32.98	32.82	32.16

Conflict index min.	335 × 335 (E.Coli)		515 × 515 (Maize)	
Row-Epitaxial. Q	5 K	10 K	5 K	10 K
Time (Pl. + Re-emb.)	930 + 1169	1732 + 1167	4082 + 4424	7856 + 4415
ACI (Pl.)	584.92	544.93	604.04	574.68
ACI (Pl. + Re-emb.)	544.23	514.10	554.87	532.74
Greedy+. Q	200	300	200	300
Time (Pl. + Re-emb.)	522 + 788	685 + 788	2131 + 2926	2757 + 2930
ACI (Pl.)	462.52	450.15	459.38	446.76
ACI (Pl. + Re-emb.)	458.02	445.98	454.84	442.55

Second, re-embedding runs very quickly for border length minimization, even on the larger chip. For conflict index minimization, the time for the re-embedding phase exceeds the time for the placement phase for both algorithms.

Finally, Greedy+ always produces better layouts in the same amount of time (or less) while looking at fewer probe candidates. In particular, for conflict index minimization on the 515×515 chip with $Q = 5\,000$ resp. 200, Greedy+ and Sequential improve the average conflict index by 18% (from 554.87 to 454.84) and need only 60% of the time, compared to Row-Epitaxial and Sequential.

Chapter 8

Analysis of Affymetrix Microarrays

General physical structure of GeneChip arrays. Control and special probes, checkerboard patterns on the borders, empty spots.

Probe pairs: perfect match (PM) and mismatch (MM) probes. Rows of PM and MM probes on the chip.

As we mentioned in Chapter 8, all GeneChip arrays that we know of can be asynchronously synthesized in 74 steps with the standard Affymetrix deposition sequence (18.5 cycles of TGCA).

This suggests that only sub-sequences of this sequence can be used as probes on Affymetrix chips. Rahmann (2006) shows that this covers about 98.45% of all 25-mers, however, it seems that Affymetrix uses an even more restrictive probe selection criterion. This is because GeneChip probes always appear in pairs, with the perfect match (PM) and the mismatch (MM) probes being located next to each other (in alternating rows of PM and MM probes), and there is evidence that the embeddings of these probes are aligned in such a way that only the middle bases are not aligned.

Chapter 9

Shortest Common Supersequence

Chapter 10

Discussion

This thesis makes several contributions to the field of...

We introduce the *conflict index*...

In Chapter X we present a algorithm...

We are not aware of any previous work that combines placement and re-embedding, even our results showed that such an approach has obvious benefits...

We have surveyed algorithms for the microarray layout problem (MLP), divided into placement, (re-)embedding, and partitioning algorithms. Because of the super-exponential number of possible layouts and the relation to the quadratic assignment problem (QAP), we cannot expect to find optimal solutions. Indeed, the algorithms we present are heuristics with an emphasis on good scalability and ideally a user-controllable trade-off between running time and solution quality, albeit without any known provable guarantees.

We believe that solution quality is more important than the running time of a placement algorithm. Even if an algorithm takes a couple of days to complete, it is time well spent given that commercial microarrays are likely to be produced in large quantities. This is specially true when we consider the time required for the whole design process of a microarray chip. Even customer designed chips, that usually have a limited number of produced units, are likely to benefit from a few extra hours of computing time.

Among the presented approaches, two recent ones (Pivot Partitioning and Greedy+) indicate that the traditional “place first and then re-embed” approach can be improved upon by merging the partitioning/placement and (re-)embedding phases. Ongoing work will show the full potential of such combined approaches.

As a suggestion for further work, we note the needed for improving the selection of probe candidates considered for filling each spot. For example, instead using a sorted list of probes, one could use a TSP tour like the one described in Sec. 3.2. However, it

is not clear if the more time-consuming TSP approach will pay off (instead, we could use this extra time to look at more candidates).

An alternative that sounds interesting would be to build some kind of “clustering” of the probes, perhaps based on a graph or a tree, in such a way that we can find similar probes more easily and spend time on candidates that are more likely to produce less conflicts.

Question: why PM and MM probes are placed together and aligned except for the middle base?

Bibliography

- W. Adams, M. Guignard, P. Hahn, and W. Hightower. A level-2 reformulation-linearization technique bound for the quadratic assignment problem. *European Journal of Operational Research*, To appear.
- H. Binder and S. Preibisch. Specific and nonspecific hybridization of oligonucleotide probes on microarrays. *Biophys J*, 89(1):337–352, Jul 2005. doi: 10.1529/biophysj.104.055343. URL <http://dx.doi.org/10.1529/biophysj.104.055343>.
- E. Çela. *The Quadratic Assignment Problem: Theory and Algorithms*. Kluwer Academic Publishers, 1997.
- P. Chase. Subsequence numbers and logarithmic concavity. *Discrete Mathematics*, 16: 123–140, 1976.
- S. A. de Carvalho Jr. and S. Rahmann. Searching for the shortest common supersequence. Technical Report 2005-03, Technische Fakultät der Universität Bielefeld, Apr 2005.
- S. A. de Carvalho Jr. and S. Rahmann. Improving the layout of oligonucleotide microarrays: Pivot Partitioning. In P. Bucher and B. Moret, editors, *Proceedings of the 6th Workshop of Algorithms in Bioinformatics*, volume 4175 of *Lecture Notes in Computer Science*, pages 321–332. Springer, 2006a. doi: 10.1007/11851561. URL <http://www.springerlink.com/content/h9r4n673058032xm>.
- S. A. de Carvalho Jr. and S. Rahmann. Microarray layout as a quadratic assignment problem. In D. Huson, O. Kohlbacher, A. Lupas, K. Nieselt, and A. Zell, editors, *Proceedings of the German Conference on Bioinformatics*, volume P-83 of *Lecture Notes in Informatics (LNI)*, pages 11–20. Gesellschaft für Informatik, 2006b.
- S. A. de Carvalho Jr. and S. Rahmann. Modeling and optimizing oligonucleotide microarray layout. In I. Mandoiu and A. Zelikowsky, editors, *Bioinformatics Algorithms: Techniques and Applications*, Wiley Book Series on Bioinformatics. Wiley, 2007. To appear.
- W. Feldman and P. Pevzner. Gray code masks for sequencing by hybridization. *Genomics*, 23(1):233–235, 1994. doi: 10.1006/geno.1994.1482. URL <http://dx.doi.org/10.1006/geno.1994.1482>.

- T. A. Feo and M. G. C. Resende. Greedy randomized adaptive search procedures. *Journal of Global Optimization*, 6:109–133, 1995.
- S. P. Fodor, J. L. Read, M. C. Pirrung, L. Stryer, A. T. Lu, and D. Solas. Light-directed, spatially addressable parallel chemical synthesis. *Science*, 251(4995):767–773, 1991.
- S. P. Fodor, R. P. Rava, X. C. Huang, A. C. Pease, C. P. Holmes, and C. L. Adams. Multiplexed biochemical assays with biological chips. *Nature*, 364(6437):555–556, Aug 1993. doi: 10.1038/364555a0. URL <http://dx.doi.org/10.1038/364555a0>.
- H. Gabow. An efficient implementation of Edmond’s algorithm for maximum matching on graphs. *J. ACM*, 23:221–234, 1976.
- J. L. Gross and J. Yellen, editors. *Handbook of graph theory*. CRC Press, 2004.
- P. Hahn, T. Grant, and N. Hall. A branch-and-bound algorithm for the quadratic assignment problem based on the hungarian method. *European Journal of Operational Research*, 108:629–640(12), 1998. doi: doi:10.1016/S0377-2217(97)00063-5. URL <http://www.ingentaconnect.com/content/els/03772217/1998/00000108/00000003/art00063>.
- S. Hannenhalli, E. Hubell, R. Lipshutz, and P. A. Pevzner. Combinatorial algorithms for design of DNA arrays. *Adv Biochem Eng Biotechnol*, 77:1–19, 2002.
- E. A. Hubbell, M. S. Morris, and J. L. Winkler. Computer-aided engineering system for design of sequence arrays and lithographic masks. United States Patent number 5,856,101, Jan 1999.
- A. Kahng, I. Mandoiu, P. Pevzner, S. Reda, and A. Zelikovsky. Border length minimization in DNA array design. In R. Guigó and D. Gusfield, editors, *Algorithms in Bioinformatics (Proceedings of WABI)*, volume 2452 of *Lecture Notes in Computer Science*, pages 435–448. Springer, 2002. URL <http://www.springerlink.com/content/pqp7c7emyk7gmx3u>.
- A. B. Kahng, I. Mandoiu, P. Pevzner, S. Reda, and A. Zelikovsky. Engineering a scalable placement heuristic for DNA probe arrays. In *Proceedings of the seventh annual international conference on research in computational molecular biology (RECOMB)*, pages 148–156. ACM Press, 2003a. doi: <http://doi.acm.org/10.1145/640075.640095>.
- A. B. Kahng, I. Mandoiu, S. Reda, X. Xu, and A. Z. Zelikovsky. Evaluation of placement techniques for DNA probe array layout. In *Proceedings of the 2003 IEEE/ACM international conference on Computer-aided design (ICCAD)*, pages 262–269. IEEE Computer Society, 2003b. doi: <http://dx.doi.org/10.1109/ICCAD.2003.65>.

-
- T. C. Koopmans and M. J. Beckmann. Assignment problems and the location of economic activities. *Econometrica*, 25:53–76, 1957.
- D. L. Kreher and D. R. Stinson. *Combinatorial Algorithms: Generation, Enumeration and Search*. CRC Press, 1999.
- Y. Li, P. Pardalos, and M. Resende. A greedy randomized adaptive search procedure for the quadratic assignment problem. In P. M. Pardalos and H. Wolkowicz, editors, *Quadratic assignment and related problems*, volume 16 of *DIMACS Series on Discrete Mathematics and Theoretical Computer Science*, pages 237–261. American Mathematical Society, 1994. URL <http://www.research.att.com/~mgcr/doc/grpqap.ps.Z>.
- S. Lin and B. W. Kernighan. An effective heuristic algorithm for the traveling salesman problem. *Operations Research*, 21:498–516, 1973.
- C. A. S. Oliveira, P. M. Pardalos, and M. G. C. Resende. GRASP with path-relinking for the quadratic assignment problem. In C. C. Ribeiro and S. L. Martins, editors, *Proc. of Third Workshop on Efficient and Experimental Algorithms (WEA04)*, volume 3059 of *Lecture Notes in Computer Science*, pages 356–368. Springer-Verlag, 2004.
- S. Rahmann. The shortest common supersequence problem in a microarray production setting. *Bioinformatics*, 19(Suppl 2):ii156–ii161, Oct 2003.
- S. Rahmann. Subsequence combinatorics and applications to microarray production, DNA sequencing and chaining algorithms. In M. Lewenstein and G. Valiente, editors, *Combinatorial Pattern Matching (CPM)*, volume 4009 of *LNCS*, pages 153–164, 2006.
- C. Savage. A survey of combinatorial Gray codes. *SIAM Review*, 39(4):605–629, 1997.
- S. Singh-Gasson, R. D. Green, Y. Yue, C. Nelson, F. Blattner, M. R. Sussman, and F. Cerrina. Maskless fabrication of light-directed oligonucleotide microarrays using a digital micromirror array. *Nat Biotechnol*, 17(10):974–978, Oct 1999. doi: 10.1038/13664. URL <http://dx.doi.org/10.1038/13664>.